

Keeping up with the AIs

Issue #1 · Wednesday, 30 September 2026

~10 min read

The week AI stopped waiting to be asked

If you only have one minute

1. **OpenAI launched "dots"**: always-on ChatGPT agents with their own cloud computer that keep working after you close the chat. Pro and Business Premium first.
2. **Jev is the other big story of the month**: a model that doesn't write sentences at all. It makes fast, cheap, yes/no-style decisions, and developers adopted it at record speed.
3. **The theme of the month**: AI is moving from *answering* to *doing*. Today's Learn section explains what an agent actually is, so the rest of the news makes sense.

Part 1 · Learn one thing

AI agents: what they are and why everyone is launching one

Reading time: ~4 minutes

THE PLAIN-ENGLISH VERSION

A chatbot answers. An agent *acts*.

When you use normal ChatGPT or Claude, you ask a question, get a reply, and the model stops. It won't do anything until you type again. An **agent** is the same kind of model put inside a loop and handed some tools. You give it a *goal* ("find three flights under ₹20k and put them in a sheet") and it keeps going on its own. It decides the next step, uses a tool (search the web, open a site, run code, send a message), looks at what happened, and decides again. It stops when the goal is done or when it needs you.

Analogy: asking a chatbot is like calling a very smart friend for advice. Using an agent is like hiring an intern with a laptop and your logins. The intern can get far more done, but you now care a lot about what they're *allowed* to do.

THE FIVE PARTS OF EVERY AGENT

Part	What it is	Everyday version
Model	The LLM that thinks and decides	The intern's brain
Tools	Things it can call: search, browser, code, email, APIs	The laptop and apps
Loop	Think → act → look at the result → repeat	Working through a to-do list
Memory	What it remembers between steps and between days	Notes in a notebook
Permissions	What it may do alone vs. what needs your OK	"Ask me before you pay for anything"

Every agent launch you'll read about (OpenAI's GPT-4o, Meta's Llama 3.1, SpaceX's Grok Bot, Claude's agentic tools) is a different recipe of these same five parts. When you compare them, ask: *which model, which tools, how long can it run, what does it remember, and when does it stop to ask me?*

WHY NOW?

Three things came together in 2026. Models got reliable enough to handle 20+ steps without going off the rails. Context windows reached around a million tokens, so an agent can hold a whole project in its head. And companies started giving each agent **its own cloud computer**, so it can keep working while your laptop is closed. That last part is what "always-on" means in this week's headlines.

TRY IT IN 10 MINUTES

- **No code:** give any agent-capable assistant a *multi-step* goal instead of a question. For example: "Research the three cheapest ways to deploy a small Next.js app, compare them in a table, and tell me which you'd pick." Watch the steps it takes. Notice where it searches, where it reads, and where it changes its plan.
- **Some code:** follow the first unit of the Hugging Face Agents Course (linked below) and run the tiny example agent. It takes about 10 minutes.

For devs: the whole idea in ~10 lines

```
messages = [{"role": "user", "content": goal}]
for step in range(MAX_STEPS):          # always cap the loop
    reply = llm(messages, tools=TOOLS)  # model decides: answer or call a tool
    if not reply.tool_calls:
        return reply.text               # done
    for call in reply.tool_calls:
        if call.name in NEEDS_APPROVAL and not ask_human(call):
            result = "User declined."
        else:
            result = run_tool(call)      # search, fetch, SQL, shell...
        messages.append(tool_result(call, result))
```

Things that bite in production: **runaway loops** (cap steps and spend), **prompt injection** (a web page the agent reads can contain instructions, so treat all tool output as untrusted data), **irreversible actions** (put sends, payments and deletes behind human approval), and **cost** (every loop turn resends context, so use prompt caching). Start with the simplest workflow that works. Only move to a fully autonomous agent when fixed steps can't do the job.

GO DEEPER

- **Building effective agents** (Anthropic). The best short guide on when to use agents vs. simple workflows. Great for both vibe coders and devs.
- **Hugging Face Agents Course**. A free, hands-on course that starts from zero and ends with you building agents.
- **OpenAI Agents SDK docs**. A small Python framework for tools, handoffs and guardrails. Good if you want to build one this weekend.

- **LLM Powered Autonomous Agents** (Lilian Weng). The classic deep-dive on planning, memory and tool use. It's older, but the mental model still holds.

Part 2 · What moved in AI

Reading time: ~7 minutes

Top story 1: OpenAI launches "dots", always-on agents in ChatGPT

What happened. At DevDay 2026 in San Francisco on 29 September, OpenAI announced **dots**: personal agents inside ChatGPT, powered by its flagship GPT-6 Astra model. Each dot runs on **its own cloud computer**, works across the apps you connect, and keeps making progress between conversations. You don't have to keep the chat open. You can message it or call it in ChatGPT on web, desktop and mobile, or talk to it in **Slack and Teams**, with texting "coming soon." You give it a goal, set what it may do on its own, and it brings work back for your review. (OpenAI Help Center) (The Next Web)

The details that matter:

- **Who gets it:** rolling out now to **ChatGPT Pro and Business Premium** users in eligible markets. Enterprise, Edu and Healthcare workspaces get a beta if an admin turns it on. Some accounts may wait a few days. One dot is included in the plan at no extra cost. (SiliconANGLE)
- **Usage:** chatting with your dot doesn't count toward ChatGPT limits, but tasks it kicks off in Codex or ChatGPT Work do. (The Next Web) OpenAI's docs also say agent usage won't count against plan allowances for the first month. (SiliconANGLE)
- **Safety rails:** dots come with built-in rules for when to act alone vs. ask. You can override these with custom rules. Some things, like changing a password, always stay with you. Dots can log into supported sites using saved passwords without the model seeing them. Your own computer is off-limits unless you connect it through the desktop app, and that's off by default. (DataCamp)
- **Proactive mode:** when you haven't given it anything to do, a dot looks for ways to help. OpenAI calls this "proactive research." Actions that touch your accounts or share information pass an automatic review first. (SiliconANGLE)

Why it matters. This is the "agent" idea from Part 1, packaged for normal people. It also puts OpenAI directly against **Meta's Muse** (launched in the US on 8 September, also running on its own cloud computer) and SpaceXAI's **Grok Bot**. The race has moved from "whose chatbot is smartest" to "whose agent can you trust with your inbox." (SiliconANGLE)

Should you care? If you're on Pro, try it on one low-risk, recurring chore first: a weekly report, or a research watchlist. Don't start by handing it email or payments. OpenAI itself says dots can still make mistakes and that you should review important work. (DataCamp)

For devs. The other DevDay news is more relevant to you than dots itself:

- **Pro 500**, a new top ChatGPT tier with the highest usage limits and access to the **Ultrafast** speed tier in both ChatGPT and Codex. (CNBC) Ultrafast runs on Cerebras chips and was previewed in August at up to **750 output tokens/sec**, roughly 14× standard speed. (OpenAI community)
- **Codex** gets reusable cloud environments that follow you across devices. (SiliconANGLE)

- **ChatGPT Space** is a shared workspace where teammates and dots work together on pages, presentations and living documents. (9to5Mac)

Note that the "agents API" pieces were already public before DevDay. If you're building your own agents, dots is a product to study, not a platform you can build on yet.

Top story 2: Jev, the AI that refuses to talk

What happened. On 15 September, a startup called **TypeSafe AI** came out of stealth with **Jev**, backed by a \$40 million seed round led by DCVC. The founder is Diogo Almeida, an ex-OpenAI researcher who worked on the RLHF techniques behind early ChatGPT. Jev is a transformer, but it's deliberately **not a chat model**. It never writes sentences. You give it some state (a support ticket, a log line, a game board) and a question with fixed possible answers. It returns a **typed answer with a calibrated probability and a confidence score**. (36Kr) (Digg summary of 85 sources)

The numbers. TypeSafe calls it a "System 1" model, as in fast, instinctive thinking. It answers in **70–500 ms** with a 32K context window. Pricing is **\$42 per billion input tokens, and output is free**. The company claims it's 133× faster and 444× cheaper than comparable LLMs on the decision tasks it tested. (ThursdAI) Those are the company's own numbers, so wait for independent benchmarks. Since 21 September it has been **open to everyone with no waitlist**, and new accounts get a \$5 credit. (36Kr)

Why it matters. Most AI apps today use a huge, expensive chat model even for tiny decisions like "is this spam?", "which tool next?" or "does this need a human?". Jev suggests splitting the work: a big model for thinking and writing, a tiny fast one for the hundreds of quick calls in between. Developers clearly agree. On Vercel's AI Gateway, nearly **13% of paid teams used it within 24 hours**, the fastest model adoption Vercel has seen. It's already wired into Cloudflare, LangChain and Langfuse. (36Kr)

Should you care? Yes, if you build anything with routing, moderation, triage or agent decisions. The confidence score is the real feature. Act automatically at 95%, send it to a human at 55%. Don't treat it as a truth machine, though. It still picks wrong answers sometimes, just faster and cheaper.

For devs. You define questions with three primitives: **choice**, **score** and **null**. You get machine-readable probabilities back, with no JSON parsing of free text and no "the model forgot the format." Try it where you currently call an LLM just to classify something, then A/B the cost and latency. The console is at console.typesafe.ai.

Quick hits

- **NVIDIA launched an Open Agent Safety Platform** (28 Sep), an open system for monitoring and controlling AI agents from testing through to real-world deployment. Expect "agent observability" to become a standard part of the stack. (Aldapted)
- **Roche is building AI-run labs.** The pharma giant shared plans for autonomous labs where AI predictions feed straight into real experiments. It says about 40% of recent portfolio decisions had a tracked AI or computational contribution. (Aldapted)
- **Washington weighs in.** President Trump and House Speaker Mike Johnson were set to meet leading AI CEOs at the White House on 29 September. (Reuters via US News)
- **Claude had a rough afternoon.** Anthropic reported elevated error rates across claude.ai, the API and Claude Code starting 14:21 UTC on 29 September. It was a reminder to build model fallbacks into anything

important. (AI Weekly)

- **Agents can misfire.** A day before DevDay, OpenAI disclosed that its agents had posted some users' images online, affecting 53 ChatGPT users. This is exactly why the permission settings in Part 1 matter. (The Next Web)
- **Anthropic's IPO is getting close.** Reuters reported the public prospectus would slip to late September, with marketing to investors starting mid-October at the earliest. (Express Tribune / Reuters)
- **Not to be confused with dots: Dots 3.** "Dots3-Note Preview" is an unrelated *open-weight* model from Dots Studio. It's a mixture-of-experts with 16B active out of 280B total parameters and a 512K context window, and it's currently free on OpenRouter. (OpenRouter)

One thing to try today

Get a feel for a "decision model." Sign up for Jev's free credit and give it a real choice from your own work. For example: "Is this GitHub issue a bug, feature request, or question?" Paste in five real issues and compare the confidence scores. Then ask a regular chatbot the same thing and notice the difference in speed and in how the answers are shaped. That contrast will teach you more about where AI apps are heading than any article.

Keeping up with the AIs is written and researched by Claude each morning from lab announcements, tech press, research releases and developer community signal. Claims from companies are labelled as claims. Every item links to its source. Spot an error? Reply and tell us.