

Keeping up with the AIs

The week AI stopped waiting to be asked

Issue #1

Wed, 30 Sep 2026

☕ ~9 min read

⚡ If you only have one minute

1. **OpenAI launched "dots"**: always-on ChatGPT agents with their own cloud computer. They keep working after you close the chat.
2. **Jev is the dark horse of the month**: a model that *never writes a sentence*. It just makes fast, dirt-cheap decisions, and developers adopted it at record speed.
3. **The big shift**: AI is moving from *answering* to *doing*. Part 1 explains what an agent really is, so the rest of the news clicks.

September in three numbers

13%

of Vercel's paid teams tried **Jev** within 24 hours of it going live there

750

tokens per second, the top speed of OpenAI's new **Ultrafast** tier

\$0.042

per *million* input tokens on Jev. Output tokens are **free**

PART 1

Learn one thing



AI agents: what they are and why everyone is launching one


4 MIN READ

MODULE 3 · BUILDING WITH AI

A chatbot answers. An agent *acts*.

Normal ChatGPT or Claude works like a vending machine: you put in a question, you get an answer, and it stops. An **agent** is the same kind of model put inside a *loop* and handed *tools*. You give it a **goal**, like "find three flights to Goa under ₹20k and put them in a sheet." It then keeps going on its own. It decides a step, does it, checks the result, and decides again, until the job is done or it needs you.

The agent loop, the one idea behind every agent launch



A chatbot does this once. An agent goes round this loop 5, 20 or 100 times until the goal is met.

Analogy: a chatbot is a smart friend you *call for advice*. An agent is an intern you *hand your laptop and logins to*. The intern gets far more done, but now it matters a lot what they're **allowed** to do.

THE FIVE PARTS OF EVERY AGENT

Part	What it is	Intern version
Model	The LLM that thinks and decides	Their brain
Tools	Search, browser, code, email, APIs	Laptop + apps
Loop	Think → act → observe → repeat	Working a to-do list
Memory	What it remembers across steps and days	Their notebook
Permissions	What it may do alone vs. ask first	"Ask before you pay for anything"

OpenAI's dots, Meta's Muse, SpaceXAI's Grok Bot and Claude's agentic tools are all different recipes built from these same five parts. **To compare any two, ask:** which model? which tools? how long can it run? what does it remember? when does it stop and ask me?

Jargon buster

- **Tool call:** the model saying "run this function with these inputs." Your code runs it and hands back the result.
- **Context window:** how much text the model can "see" at once. Some models now handle around a million tokens, roughly a thick novel.
- **Prompt injection:** hidden instructions in a web page or file that try to hijack the agent reading it.
- **Human-in-the-loop:** the agent pauses and asks you before doing something important.

WHY NOW?

Three things came together in 2026. Models got reliable enough to handle **20+ steps** without wandering off. Context windows reached around **a million tokens**, so an agent can hold a whole project in its head. And companies started giving each agent **its own cloud computer**, so it keeps working while your laptop is closed. That last part is what "always-on" means in this week's headlines.

Try it in 10 minutes

- **No code:** give any agent-capable assistant a *multi-step goal* instead of a question. For example: "Compare the three cheapest ways to deploy a small Next.js app, put them in a table, and pick one." Watch where it searches, where it reads, and where it changes its plan.
- **Some code:** run the tiny example agent from Unit 1 of the Hugging Face Agents Course (linked below).

For devs: the whole idea in ~10 lines

```
messages = [{"role": "user", "content": goal}]
for step in range(MAX_STEPS):          # always cap the loop
    reply = llm(messages, tools=TOOLS)  # model decides: answer or call a tool
    if not reply.tool_calls:
        return reply.text               # done
    for call in reply.tool_calls:
        if call.name in NEEDS_APPROVAL and not ask_human(call):
            result = "User declined."
        else:
            result = run_tool(call)      # search, fetch, SQL, shell...
        messages.append(tool_result(call, result))
```


What bites in production: *runaway loops* (cap steps and spend), *prompt injection* (treat every tool output as untrusted data), *irreversible actions* (put sends, payments and deletes behind approval), and *cost* (every turn resends context, so use prompt caching). Start with the simplest fixed workflow that works. Only go fully autonomous when fixed steps can't do the job.

GO DEEPER

- **Building effective agents** (Anthropic). The best short read on agents vs. simple workflows.
- **Hugging Face Agents Course**. Free and hands-on, from zero to building your own.
- **OpenAI Agents SDK docs**. A small Python framework. Good weekend project material.
- **LLM Powered Autonomous Agents** (Lilian Weng). The classic deep-dive on planning, memory and tools.

7 min read · Everything since the start of the month that you need to know, with the last 24 hours up top.

September 2026: the month agents went mainstream




Top story 1 · OpenAI launches "dots", always-on agents in ChatGPT

BIG DEAL

PRO & BUSINESS PREMIUM FIRST

What happened. At DevDay 2026 in San Francisco on 29 September, OpenAI announced **dots**: personal agents inside ChatGPT, powered by its flagship GPT-6 Astra model. Each dot has **its own cloud computer**, works across the apps you connect, and keeps making progress *between* conversations. (OpenAI Help Center) You can message it or call it in ChatGPT, or talk to it in **Slack and Teams**. Texting is "coming soon." (The Next Web)

How a dot works



Your own computer stays out of reach unless you connect it through the desktop app. That's off by default. Sources are in the story.

The details that matter:

- **Who gets it:** ChatGPT Pro and Business Premium users in eligible markets, rolling out now, so some accounts may wait a few days. **One dot is included** at no extra cost. (SiliconANGLE) Enterprise, Edu and Healthcare get a beta once an admin enables it. (OpenAI Help Center)
- **Usage:** chatting with your dot doesn't count toward ChatGPT limits, but tasks it starts in Codex or ChatGPT Work do. (The Next Web)

- **Safety rails:** dots have built-in rules for acting alone vs. asking, and you can override them with custom rules. They can log into sites with saved passwords **without the model ever seeing them.** (DataCamp)
- **Proactive mode:** with nothing to do, a dot goes looking for ways to help. OpenAI calls this "proactive research." (SiliconANGLE)

Why it matters. This is Part 1's agent, packaged for normal people. It also sets OpenAI directly against **Meta's Muse**, which launched in the US on 8 September and also runs on its own cloud computer. The race is no longer "whose chatbot is smartest." It's "**whose agent do you trust with your inbox?**" (SiliconANGLE)

WHO SHOULD CARE



Everyone



Vibe coders



Developers




Should you care? If you're on Pro, start your dot on **one low-risk, recurring chore**, like a weekly report or a research watchlist. Don't start with email or payments. OpenAI itself says dots can still make mistakes, so review anything important. (DataCamp)

For devs: the rest of DevDay matters more to you

- **Pro 500** is a new top ChatGPT tier with the highest limits plus the **Ultrafast** speed tier in ChatGPT and Codex. (CNBC)
- **Codex** gets reusable cloud environments that follow you across devices. (SiliconANGLE)
- **ChatGPT Space** is a shared workspace where teammates and dots edit pages, decks and living docs together. (9to5Mac)

What "Ultrafast" feels like: time to write a 500-token answer

Standard speed



~9 sec

Ultrafast (Cerebras chips)



~0.7 sec

Ultrafast was previewed in August at up to 750 output tokens/sec, "up to 14x" standard (OpenAI community). The Standard figure is our estimate derived from that claim, not an OpenAI number.

Top story 2 · Jev, the AI that refuses to talk

DARK HORSE

FREE TO TRY

What happened. On 15 September, a startup called **TypeSafe AI** came out of stealth with **Jev** and a \$40M seed round led by DCVC. The founder is Diogo Almeida, an ex-OpenAI researcher who worked on the techniques behind early ChatGPT. Jev is **deliberately not a chat model**. You give it some state (a support ticket, a log line, a game board) and a question with fixed answer options. It sends back a **pick, a probability and a confidence score**. No sentences. (36Kr) (Digg, 85 sources)

Same question, two kinds of AI

💬 Chat model (LLM)

- Writes an answer word by word
- Great for writing, reasoning, code
- Overkill (slow and pricey) for tiny yes/no calls

"This looks like a bug report, since the user mentions a crash after..."

VS

🎯 Decision model (Jev)

- Returns a **typed pick + probability**
- Answers in **70–500 ms**
- Output tokens are **free**

bug	0.91
feature	0.06
question	0.03

The Jev output shown is an illustration of the format, not a real run. Speed and pricing are from TypeSafe via ThursdAI.

The numbers. TypeSafe calls Jev a "**System 1**" model, as in fast, gut-feel thinking. It has a 32K context window and costs **\$42 per billion input tokens**, with output free. The company claims it's 133× faster and 444× cheaper than comparable LLMs on the decision tasks it tested. These are **the company's own numbers**, so wait for independent benchmarks. (ThursdAI) Since 21 September it has been **open to everyone** with a \$5 starter credit. (36Kr)

Why it matters. Most AI apps use a huge, expensive model even for tiny calls like "is this spam?" or "which tool next?". Jev suggests a split: **a big model for thinking and writing, a tiny fast one for the hundreds of quick decisions in between**. Developers clearly liked the idea. Nearly 13% of Vercel's paid AI Gateway teams tried it within 24 hours, and it's already wired into Cloudflare, LangChain and Langfuse. (36Kr)

WHO SHOULD CARE



Vibe coders



Developers



For devs. You define questions with three primitives, **choice**, **score** and **null**, and get machine-readable probabilities back. There's no free text to parse. The confidence score is the real feature: auto-act at 95%, send to a human at 55%. Try it anywhere you currently call an LLM just to classify something, and A/B the latency and cost.

⚡ Quick hits

🛡️ NVIDIA's Open Agent Safety Platform.

An open system for monitoring and controlling agents from testing to deployment (28 Sep). "Agent observability" is becoming a standard part of the stack.

(AIdapted)

🧪 Roche is building AI-run labs.

AI predictions will feed straight into real experiments. Roche says about 40% of recent portfolio decisions had a tracked AI or computational contribution. (AIdapted)

🏛️ Washington weighs in.

President Trump and House Speaker Mike Johnson were set to meet leading AI CEOs at the White House on 29 September. (Reuters via US News)

🔊 Claude had a rough afternoon.

Anthropic reported elevated errors across claude.ai, the API and Claude Code from 14:21 UTC on 29 Sep. If you ship on one model, build a fallback. (AI Weekly)

⚠️ Agents can misfire.

The day before DevDay, OpenAI disclosed that its agents had posted some users' images online, affecting 53 people. This is exactly why permissions matter. (The Next Web)

📈 Anthropic's IPO is getting close.

Reuters reported the public prospectus was slipping to late September, with investor marketing from mid-October at the earliest. (Express Tribune / Reuters)

🔵 Not to be confused: Dots 3.

"Dots3-Note Preview" is an unrelated *open-weight* model from Dots Studio: 16B active of 280B total parameters, a 512K context window, and free on OpenRouter for now. (OpenRouter)

📅 Coming up.

Watch for the White House meeting outcome, Anthropic's public S-1, independent Jev benchmarks, and dots expanding beyond Pro.

🎯 One thing to try today: feel what a "decision model" is

Grab Jev's free credit and give it a real choice from your own work. For example: *"Is this GitHub issue a bug, a feature request, or a question?"* Paste in five real issues and look at the confidence scores. Then ask a regular chatbot the same thing. The difference in **speed** and in the **shape of the answer** teaches more about where AI apps are heading than any article.

🧠 Pop quiz: 30 seconds, no Googling

1. What makes an agent different from a chatbot?

- ☐ A It uses a bigger model
- ☐ B It runs a loop with tools toward a goal
- ☐ C It has a voice

2. What does Jev return?

- ☐ A A paragraph
- ☐ B A typed pick with a probability
- ☐ C An image

3. Where does your dot do its work?

- ☐ A On your laptop
- ☐ B On its own cloud computer
- ☐ C Only inside Slack

Answers (no peeking): 1. B: goal + tools + loop - 2. B: decisions, not sentences - 3. B: its own cloud PC, always on

Keeping up with the AIs is researched and written by Claude every morning from lab announcements, tech press, research releases and developer chatter. Company claims are labelled as claims, and every item links to its source. Spot an error or want a topic covered? Just reply. ✉️