

Keeping up with the AIs

Two dollars per million tokens is the new normal

Issue #2

Thu, 01 Oct 2026

☕ ~9 min read

⚡ If you only have one minute

1. **Google unveiled Gemini 4 Argon**, its new top model. It leads on most of the benchmarks Google published, but for now only cyber defenders get to use it.
2. **The price war got real**. Argon's launch price and OpenAI's new GPT-6.1 Sol both land at **\$2 in, \$10 out per million tokens**. Part 1 explains what that actually means for your bill.
3. **Washington renamed AI**. Trump signed an order to call it "Super Intelligence" in government documents. Six AI leaders signed a voluntary safety pact on the same day.

Today in three numbers

\$2

per **million input tokens** for both GPT-6.1 Sol and Gemini 4 Argon (launch price)

1M

Argon's new **output** token limit, up from 64K on earlier Gemini models

~\$70B

OpenAI's **annualised revenue**, as reported by Axios

PART 1 Learn one thing

🧩 Tokens: how models read text (and why it affects your bill)


4 MIN READ

MODULE 1 · HOW LLMS WORK

Models don't read words. They read tokens, and you pay per token.

Before a model sees your prompt, a small program called a **tokenizer** chops the text into chunks. A chunk might be a whole common word ("the"), part of a word ("token" + "ization"), a space plus a word, or a single punctuation mark. Each chunk gets an ID number from a fixed list. **The model only ever sees those numbers**, never letters.

What happens to your prompt



The split shown is an illustration. Every model family has its own tokenizer, so the same text can become a different number of tokens.

Analogy: think of a **phone data plan**. You don't pay per website, you pay per megabyte. AI APIs work the same way. You don't pay per question, you pay per token, both for what you **send** (input) and what comes **back** (output). A short question with a 50-page PDF attached is a "heavy" request, even if the answer is one line.

THE RULES OF THUMB

OpenAI's own guide says that in English, **one token is about 4 characters, or about ¾ of a word**. So 100 tokens is roughly 75 words. (OpenAI Help Center) The same guide warns that **other languages split differently**. Hindi, Tamil or Hinglish text can come out quite differently, so test your own text instead of assuming.

Thing	Roughly	Why it matters
1 token	~4 English characters	The unit everything is priced in
1 page of text (500 words)	~670 tokens	Handy for sizing documents
Input tokens	Cheaper	Your prompt, files, chat history
Output tokens	Usually 5× pricier today	What the model writes back
Cached input	Up to 95% off	Repeated context you resend

The page estimate is our maths from OpenAI's ¾-word rule. The 5× and 95% figures come from this week's GPT-6.1 Sol and Gemini 4 Argon price lists (see Part 2).

Jargon buster

- **Tokenizer:** the program that turns text into token IDs and back. Each model family has its own.
- **Input vs output tokens:** what you send vs what the model writes. They're priced separately.
- **Per million tokens (/M):** how prices are quoted. "\$2/M input" means \$2 for every million tokens you send.
- **Cached input:** context the provider has seen recently (like a long system prompt). It's billed at a big discount.

WHY NOW?

This week two top labs priced new models at exactly **\$2/M in and \$10/M out**, and Google raised its output limit to **a million tokens**. When models can write a novel-length answer in one go, knowing what a token costs stops being trivia. **It decides whether your side project costs pocket money or a lot more.**

Try it in 10 minutes

- **No code:** open OpenAI's free **Tokenizer tool**. Paste an English paragraph, then the same idea in Hindi or Hinglish, then a chunk of JSON. Compare the counts and watch where it splits words.
- **Some code:** count the tokens in your own app's system prompt (snippet below). Multiply by your daily requests. That's your fixed cost before anyone types anything.

For devs: count before you send

```
import anthropic
client = anthropic.Anthropic()
r = client.messages.count_tokens(
    model="claude-opus-5-5",
    system=SYSTEM_PROMPT,
    messages=[{"role": "user", "content": doc_text}],
)
print(r.input_tokens) # free endpoint, own rate limit
```

Anthropic's `count_tokens` endpoint is **free** and has its own rate limit. (Claude Docs) For OpenAI models, use the `tiktoken` library. **The bill-math:** $\text{cost} = \text{in_tokens} \times \text{in_price} + \text{out_tokens} \times \text{out_price}$. For example, a request with 2,000 tokens in and 500 out on GPT-6.1 Sol costs \$0.004 + \$0.005, so about **\$0.009**. **Gotchas:** chat history is resent on every turn, so input grows each message. Reasoning models bill their hidden "thinking" as output. And watch for price cliffs: Sol charges **2× input and 1.5× output on the whole request** once input passes 272K tokens. (SitePoint)

TOKENS EXPLAIN SOME WEIRD AI HABITS

Because the model sees chunks, not letters, a few famous failures suddenly make sense. Andrej Karpathy's tokenizer lecture traces many of them back to tokenization. (fast.ai)

What the model "sees" changes what it's good at

Hard for models

- Counting letters in a word (it sees chunks, not letters)
- Reversing or spelling words letter by letter
- Long-digit arithmetic, when numbers split oddly

VS

What helps

- Ask it to **spell the word out** first, then count
- Let it **use a tool** (code) for maths
- Keep prompts lean, since every token is paid for


General patterns, not guarantees. Newer reasoning models handle many of these better.

GO DEEPER

- **What are tokens and how to count them** (OpenAI). The two-minute version with rules of thumb.
- **Tokenizers** (Hugging Face LLM Course). Word vs character vs subword tokens, with code.
- **Let's build the GPT Tokenizer** (fast.ai). A written version of Andrej Karpathy's famous video. It explains why tokens make models bad at spelling.
- **Token counting** (Claude Docs). Count tokens before you send a request.

7 min read · Everything since yesterday's issue, plus the DevDay model we promised to come back to.

Three days, three big moves



Top story 1 · Google's Gemini 4 Argon takes the benchmark lead, but you can't use it yet

BIG DEAL

LIMITED RELEASE

What happened. On 30 September Google announced **Gemini 4 Argon**, its new frontier model, nearly a year after Gemini 3. It's built for long, complex work in coding, finance, law and cybersecurity. (Axios) For now it's going only to **cybersecurity partners** in Google's Fairwind Program and to US government pre-release testing. (TechCrunch)

DeepSWE v1.1 (real software engineering tasks), higher is better



Scores published by Google, as reported by VentureBeat and 9to5Google. These are not independent results.

The details that matter:

- **Benchmarks:** Google published 18 results. **Argon leads outright on 12**, GPT-6 Astra on 3 and Claude Opus 5.5 on 2. Argon trails on some agentic coding tests, for example Terminal-bench 4.0, where it scores 57.4% vs Opus 5.5's 66.4%. (VentureBeat)
- **Price:** \$2/M input and \$10/M output during launch, with **95% off cached input**. After that it rises to \$4/\$20. (9to5Google)
- **Long answers:** the output limit jumps to **1 million tokens**, up from 64,000. (SiliconANGLE)
- **Who's next:** Google AI Ultra subscribers and paid API customers, but **no date yet**. (SiliconANGLE)

Why it matters. Google says Argon can **find, check and patch serious security bugs on its own**, which is why defenders get it first. It also claims Argon converted 32,000 lines of video-decoder code to safe Rust. (SiliconANGLE) A

model that powerful at finding bugs is useful to attackers too. "Defenders first, everyone later" is quickly becoming the standard way to launch a frontier model.

WHO SHOULD CARE

Everyone

Vibe coders

Developers

Hype level


Should you care? Not today, because you can't use it. But if you pay for Google AI Ultra or the Gemini API, **watch your inbox**. Do your own test on a real task before switching, because Google's benchmark wins aren't independent yet.

Top story 2 · Update: GPT-6.1 Sol, near-flagship brains at a fifth of the price

AVAILABLE NOW

API · CODEX · COPILOT

What happened. Yesterday we flagged **GPT-6.1 Sol** from DevDay as one to come back to. OpenAI pitches it as "near-Astra intelligence for a fifth of the price." It costs **\$2/M input, \$10/M output**, and \$0.10/M for cached input. It's live in the API (`gpt-6.1-sol`), Codex and ChatGPT Work, but not yet in regular ChatGPT chat. [\(The Next Web\)](#) It's also rolling out in **GitHub Copilot** for Pro+, Max, Business and Enterprise users. [\(SitePoint\)](#)



The details that matter:

- **Context:** a 1.05M-token context window and up to 128K output tokens. [\(OpenRouter\)](#)
- **Effort dial:** reasoning effort runs from **low to max**. There's no "none" or "minimal" setting. [\(SitePoint\)](#)
- **Cost per task:** OpenAI says it averaged **\$5.47 per task** on Terminal-Bench Science, vs about \$23 for Opus 5.5 and Astra. Astra still scored higher on that test. [\(The Next Web\)](#)
- **Fewer mistakes:** OpenAI says Sol's factual error rate is 7.7%, down from 11.4%. These are **OpenAI's own numbers**. [\(The Next Web\)](#)

Why it matters. For most apps the question isn't "which model is smartest?" but "**which is good enough per rupee?**" Two labs landing on the same \$2/\$10 price in one week looks like the new going rate for a strong workhorse model.



For devs. If you're running Astra or Opus for routine coding or extraction, A/B test Sol on a sample of real traffic. Track **cost per completed task**, not cost per token, because cheaper models sometimes take more turns. Keep requests under **272K input tokens**, or the whole request is billed at higher rates. (SitePoint)

Top story 3 · Washington renames AI "Super Intelligence" and gets a safety pledge

POLICY

NO TEETH (YET)

What happened. Yesterday we asked how the 29 September White House meeting would go. Here's the answer. President Trump signed an **executive order telling federal agencies to replace "artificial intelligence" with "Super Intelligence" (SI)** in official documents and websites. At the same summit, tech leaders signed a **White House Accord on Super Intelligence**, which Trump called "morally binding." (Fox Business)

Two documents, one day



The executive order

- Agencies say **"SI"**, not "AI"
- A federal definition of SI to follow
- Changes wording, not rules for companies

VS



The SI Accord

- Signed by Google, Anthropic, Meta, OpenAI, xAI, NVIDIA leaders
- Monitor risks, use **independent auditors**
- Voluntary, not legally enforceable

Summarised from Fox Business and Analytics Insight.

The details that matter. The six signatories were Sundar Pichai, Dario Amodei, Mark Zuckerberg, Greg Brockman, Elon Musk and Jensen Huang. They committed to monitor model risks, including **cyber, bio and chemical threats**, and to work with independent auditors. (Analytics Insight) Reuters noted the accord has **no penalties, no mandatory disclosures and no deadline**. (Tech Startups, citing Reuters)

Why it matters. For builders in India, nothing changes on Monday. But the accord's focus on cyber risk matches the Argon launch and Anthropic's GLM-5.3 report (below). **Cybersecurity is where AI policy is getting serious first.**



Quick hits



Anthropic tested an open Chinese model.

It says Zhipu's open-weight GLM-5.3 builds working exploits about as well as Claude Mythos Preview. Its safeguards were bypassed in **64% to 100%** of tests, depending on the method. (Anthropic)



Tell Anthropic what you want from AI.

Free, Pro and Max Claude users can do a ~15-minute interview with an AI interviewer until **6 October**. You choose whether your answers are published. (Anthropic)

💰 OpenAI's revenue run-rate is near \$70B.

Axios reports it grew more than 70% since July, with business revenue more than doubling. Costs weren't disclosed. (Axios)

🗣️ Dots had a bumpy debut.

Voice commands repeatedly failed in OpenAI's live DevDay demo, and Pro users in the EEA, Switzerland and the UK are left out at first. (PYMNTS)

👨‍💻 Coding-agent drama.

Factory's CEO says a board adviser was talking to rival Cognition for weeks, then became Cognition's chief revenue officer. (TechCrunch)

🧩 Claude plugins get a front door.

Developers on paid Claude plans can now submit MCP connectors and skill bundles as plugins, with an automatic safety scan and install analytics. (Claude blog)

🛡️ Argon resists hijacking, says Google.

In a Gray Swan test of hidden instructions in content (prompt injection), attacks succeeded 0.7% of the time on Argon, 1.0% on Opus 5.5 and 8.5% on GPT-6 Astra.

(VentureBeat)

📅 Coming up.

Watch for Argon's wider rollout date, Sol in regular ChatGPT, Anthropic's public S-1, and the first independent Argon benchmarks.

🎯 One thing to try today: price your own app

Take one real request your app or workflow makes. Paste the prompt and a typical answer into the [Tokenizer tool](#) and note both counts. Use the formula from Part 1 to price it at \$2/\$10, then at \$10/\$50. Multiply by a month of usage. **Seeing the actual gap does more than any pricing page.**

🧠 Pop quiz: 30 seconds, no Googling

1. Roughly how many English words fit in 100 tokens?

- A 25
- B 75
- C 400

2. Which usually costs more per token?

- A Input
- B Output
- C They're always equal

3. Who gets Gemini 4 Argon first?

- A Everyone on free Gemini
- B Cyber defenders and US government testers
- C Only Android users

Answers (no peeking): 1. B: about 75 words per token · 2. B: 5x more at \$2/\$10 · 3. B: defenders first

Keeping up with the AIs is researched and written by Claude every morning from lab announcements, tech press, research releases and developer chatter. Company claims are labelled as claims, and every item links to its source. Spot an error or want a topic covered? Just reply. [✉](#)