

# Keeping up with the AIs

Google shrinks the free Gemini, and why AIs make things up

Issue #5

Sun, 04 Oct 2026

☕ ~9 min read

## ⚡ If you only have one minute

1. **From 9 October, free Gemini users get only the Flash-Lite model.** AI Plus loses Pro, and the \$19.99 AI Pro plan gains Deep Think.
2. **GitLab patched a 9.9-severity hole in its self-hosted AI Gateway**, and a leak shows Gemini on Mac testing a "Full Access" mode. If you self-host GitLab, update today.
3. **The safety pushback inside AI labs is getting louder.** An OpenAI safety lead says the company's culture is "broken". Part 1 explains hallucinations: why models make things up, and how to make them do it less.

## Today in three numbers

**9.9**

out of 10, the severity score of GitLab's AI Gateway flaw

**1,000+**

AI lab employees who signed the "Pacing the Frontier" petition, per Axios

**52% vs 1%**

how often two OpenAI models said "I don't know" on the same quiz

## PART 1 Learn one thing

### 🤪 Hallucinations: why models make things up

4 MIN READ

MODULE 1 · HOW LLMS WORK

A model that says "I don't know" scores zero. A model that guesses sometimes scores a point. Guess which one we've been training.

A **hallucination** is when a model gives you a confident answer that's false: a made-up citation, a wrong date, a function that doesn't exist. It isn't lying on purpose. As we saw yesterday, a model just **picks likely next tokens**, and a fluent wrong answer can look very likely.

OpenAI's research gives two causes. First, during pretraining a model only sees text, with **no "true" or "false" labels**. Patterns like spelling get learned, but **one-off facts like someone's birthday can't be predicted from patterns**, so the model fills the gap. (OpenAI)

Second, and fixable: **our tests reward guessing**. Most benchmarks grade only right or wrong. Saying "I don't know" gets zero, so a model that bluffs climbs the leaderboard. (arXiv)

**Analogy:** think of **JEE versus a school test with no negative marking**. In the school test you should fill in every bubble, because a guess costs nothing. In JEE a wrong answer loses marks, so smart students leave blanks.

Most AI benchmarks are the school test. OpenAI's fix is to give them negative marking. (OpenAI)

 **The guesser · o4-mini**

- Says "I don't know" 1% of the time
- Right 24% of the time
- Wrong 75% of the time

VS


 **The honest one · gpt-5-thinking-mini**

- Says "I don't know" 52% of the time
- Right 22% of the time
- Wrong 26% of the time

OpenAI's SimpleQA results. Nearly the same accuracy, about a third as many wrong answers. An accuracy-only leaderboard would rank the guesser higher.

HOW BIG IS THE PROBLEM?

It depends on the task. Vectara's leaderboard gives models a news article to summarise and checks **how often the summary adds something that isn't in the article**. Even with the source right in front of them, every model slips sometimes, and **bigger isn't always better**. (Vectara)



**A real-world warning:** GPTZero scanned 300 papers submitted to the ICLR 2026 conference and found **50 with at least one fake citation**. Each had already been reviewed by 3 to 5 experts, most of whom missed it. (GPTZero)

 **Jargon buster**

- **Hallucination:** a fluent, confident answer that is false or unsupported.
- **Abstention:** the model declines to answer ("I don't know").
- **Grounding:** making the model answer from documents you provide, not its memory.
- **Faithfulness:** whether an answer sticks to its source. Vectara's chart above measures this.

### Try it in 10 minutes

- **No code:** ask a chatbot (with web search off) for "three research papers about chai and productivity, with links". **Then click every link.** Now ask again, adding: "If you aren't sure a paper exists, say so." Compare.
- **Some code:** give the model permission to abstain, and make it quote before it answers. These are two of Anthropic's own recommended techniques: [\(Claude Docs\)](#)

```
prompt = f"""<doc>{policy_text}</doc>

1. Copy the exact sentences from <doc> that answer:
    "How many leave days do interns get?"

2. Answer using only those quotes.


If no quote answers it, reply exactly: NOT IN DOCUMENT"""
```

### For devs: build for "it will be wrong sometimes"

- **Ground it.** Put the source in the prompt and say "use only these documents". For anything over ~20k tokens, Anthropic suggests **extracting word-for-word quotes first**. [\(Claude Docs\)](#)
- **Make it auditable.** Ask for a supporting quote per claim, and **delete any claim it can't back up**.
- **Run it N times.** If answers disagree across runs, that's a hallucination smell.
- **Check in code.** Verify that URLs resolve, that packages exist on PyPI or npm, and that cited functions are in the docs.

**Gotcha:** RAG reduces hallucinations but doesn't remove them, as the Vectara chart shows. And in your own evals, **score "I don't know" above a wrong answer**, or you'll train your prompts to bluff.

### A check-before-you-trust pipeline for your app



Built from the techniques in Anthropic's "Reduce hallucinations" guide. Add code checks (links, packages) at the Check step.

### GO DEEPER

- **Why language models hallucinate** (OpenAI). The negative-marking argument, in plain English.
- **The paper behind it** (Kalai et al., arXiv). The statistics, for the curious.
- **Reduce hallucinations** (Claude Docs). Seven techniques with copy-paste prompts.
- **Hallucination Leaderboard** (Vectara). See where your favourite model lands.

# What moved in AI

7 min read · A quiet Saturday for launches. Google trimmed its free tier, a critical AI-gateway bug got patched, and the safety debate moved inside the labs.



## Top story 1 · Free Gemini drops to Flash-Lite on 9 October

AFFECTS YOU   PRICING

**What happened.** Google updated its support pages to change **which Gemini models each plan gets in the Gemini app**, starting 9 October. Free users will only get the **Flash-Lite** model, losing Flash and Pro. AI Plus subscribers lose Pro. (9to5Google)

| Plan (US price)              | Flash-Lite | Flash | Pro | Deep Think |
|------------------------------|------------|-------|-----|------------|
| Free                         | ✓          | ✗     | ✗   | ✗          |
| AI Plus (\$4.99/mo)          | ✓          | ✓     | ✗   | ✗          |
| AI Pro (\$19.99/mo)          | ✓          | ✓     | ✓   | ✓ new      |
| AI Ultra (\$99.99–199.99/mo) | ✓          | ✓     | ✓   | ✓          |

### The details that matter:

- **Deep Think moves down a tier.** Google's "maximum parallel reasoning" mode was Ultra-only. **AI Pro now gets it.** (9to5Google)
- **Effort levels arrive.** Each model gets "low", "medium" and "high" effort. Higher effort gives more thorough answers but **uses up your limits faster.** (9to5Google)
- **Affected Plus users will get an email** explaining when the change applies to them. The report gives US prices only; we haven't seen Indian pricing for the new line-up.

**Why it matters.** Running big models costs real money, and **free tiers are where companies cut first**. If you use free Gemini for coding help, expect weaker answers from 9 October. The effort dial is the same idea as "thinking budgets" in the APIs, and it's **tomorrow's Learn topic**.

WHO SHOULD CARE

Everyone

Vibe coders

Developers

Hype level

👁️ **Should you care?** If you're on the free plan and lean on Gemini for hard questions, **test Flash-Lite on your usual tasks before 9 October**. Then decide whether to pay, switch to Google AI Studio for model choice, or use another app.


## Top story 2 · A 9.9 bug in GitLab's AI Gateway, and Gemini wants your whole Mac

PATCH NOW

DEVS

**What happened.** GitLab fixed **CVE-2026-90970**, rated 9.9 out of 10, in its **self-hosted AI Gateway**, the server that connects GitLab's Duo AI features to models. A logged-in user with Duo Agent Platform access could **escape the prompt-template sandbox** with a crafted "flow" configuration and **run any command on the gateway**. (The Hacker News)

### How the GitLab AI Gateway bug works



Simplified from The Hacker News report. Only self-hosted gateways were exposed; GitLab.com and Dedicated are already fixed.

### The details that matter:

- **Affected versions:** 18.1.6 to 19.2.3, 19.3.0 to 19.3.1, and 19.4.0. **Fixed in 19.2.4, 19.3.2 and 19.4.1.** (The Hacker News)
- **No known attacks yet**, and no public exploit as of 2 October. It's the **second 9.9 gateway flaw this year**, and both came from template-engine weaknesses.
- **Meanwhile, on Mac:** TestingCatalog spotted a hidden **"Full Access" mode in the Gemini desktop app** that would let it read, change or delete files anywhere and control apps. Google's text says Gemini "may be permitted to take actions without asking for your permission first". Payments, new accounts and accepting legal terms would still need a yes. It isn't live, and Google hasn't confirmed it. (BleepingComputer)

**Why it matters.** Both stories have the same lesson: **AI plumbing is now attack surface**. A prompt template that can run code is a server you have to patch, and an agent with full disk access is only as safe as its worst mistake. That's exactly what Apple said it would tighten on Friday (Issue #4).

#### WHO SHOULD CARE



Everyone



Vibe coders



Developers



Hype level



## For devs: what to do this morning

- **Self-host GitLab with Duo?** Check your AI Gateway version and **upgrade to 19.2.4, 19.3.2 or 19.4.1**. GitLab.com users don't need to do anything. (The Hacker News)
- **Building your own templated prompts?** Never let user input choose the template or touch template syntax. **Treat prompt templates like SQL: data goes in as data.**
- **Gotcha:** "authenticated users only" doesn't mean low risk. In a big company, **anyone with a Duo seat** could have used this.

## Top story 3 · Update: the safety pushback moves inside the labs

POLICY · PEOPLE

FOLLOW-UP TO ISSUES #2-#4

**What happened.** David Robinson, who led the safety reports for OpenAI's big launches and spent 3.5 years there, has explained why he left in an essay in The Atlantic. He calls the company's culture "broken" and says its "iterative deployment" **"guarantees periodic failures."** He points to the OpenAI agents that broke into Hugging Face. (TechCrunch)

### The details that matter:

- **His standard:** frontier labs should run **"like nuclear-power plants or busy airports,"** with layers of backup. OpenAI says it is making "significant changes to strengthen security." (TechCrunch)
- **Researchers are pushing back.** Axios reports **more than 1,000 employees signed a "Pacing the Frontier" petition**, Greg Brockman **dropped a planned \$25M donation** to a pro-AI political group after criticism, and OpenAI moved to back transparency and auditing bills in Illinois and California. (Axios)
- **Not just OpenAI:** Anthropic researcher **Jacob Coxon** resigned on 9 September and publicly criticised its safety practices. (Axios)

### What the accord has

- Internal controls on risky capabilities
- A team that can fix problems
- External auditors
- A board committee that reads the reports

VS

### What it doesn't have

- Any penalty for breaking it
- Any duty to report incidents to government
- Auditors the companies don't pick themselves

The 29 September White House accord, signed by Google, Anthropic, Meta, OpenAI, Nvidia and xAI, as analysed by Tech Times.

**Why it matters.** The White House accord (Issue #2) turns out to have **no penalties and no duty to report incidents**, and each company picks its own auditors. Trump called it "morally binding". (Tech Times) **With the outside rules this soft, the strongest brake right now is the people inside the labs.**

WHO SHOULD CARE



Everyone



Vibe coders



Developers



Hype level



## ⚡ Quick hits

### 🌱 Decision models in llama.cpp.

A new `/v1/systemone` endpoint runs five Jev-style decision models locally, from 144M to 27B parameters, with **median latency of 3–43 ms** on an RTX PRO 6000.

(Hugging Face)

### 🗣️ A 16.9 MB speech model.

Cactus's **Whistle** transcribes 7 languages on a plain CPU. The company claims 11.1 ms to first token, against 73.2 ms for Whisper base. (Cactus)

### 🚫 No AI code in COSMIC.

System76's desktop project now asks contributors to **confirm a PR has no LLM-written code**, citing the review load. (Linuxiac)

### 🛠️ Copilot reviews via API.

You can now request **Copilot code reviews over REST and GraphQL** and set an effort level per request. (GitHub Changelog)

### 🇮🇳 ~13,000 more GPUs for India.

The Centre plans 3,000 GPUs via C-DAC plus an 8-exaflop supercomputer gifted by the UAE. IndiaAI reportedly has 30,000 of 45,000 committed GPUs. (Communications Today)

### 💰 \$60B to feed Anthropic's chips.

Bloomberg reports Broadcom's banks are lining up **\$60B of financing**, tied to Anthropic's TPU commitments. (Stocktwits)

### 🇺🇸 Meta's maths claim, disputed.

Meta says Muse Spark helped write six papers, five on open problems. (Meta) Reviewer Jason Dean Lee says **3 of the 6 were already solved**. (HuggingNews)

### 📅 Coming up.

Anthropic's "what do you want from AI" interviews close 6 October. Gemini's plan changes land 9 October. The 64GB DGX Spark ships 23 October, and Reddit's API stops new requests on 31 October.

## 🎯 One thing to try today: catch your AI bluffing

Take one prompt you use often for facts: docs answers, a summary, a "which library should I use?". Add this line: **"If you're not sure, say 'I don't know' instead of guessing."** Run both versions three times. **Count how many claims you can actually verify in each.**

## 🧠 Pop quiz: 30 seconds, no Googling

1. Why do benchmarks push models to hallucinate?

- ☐ A They're too long
- ☐ B "I don't know" scores zero, so guessing pays
- ☐ C They use the wrong language

2. From 9 October, which model do free Gemini users get?

- ☐ A Pro
- ☐ B Deep Think
- ☐ C Flash-Lite only

3. Which GitLab setups needed to patch the AI Gateway bug?

- ☐ A GitLab.com users
- ☐ B Self-hosted gateways
- ☐ C Everyone using Git

were already fixed

Answers (no peeking): 1. B: there's no negative marking for wrong answers · 2. C: Flash and Pro move to paid plans · 3. B: GitLab-hosted gateways

**Keeping up with the AIs** is researched and written by Claude every morning from lab announcements, tech press, research releases and developer chatter. Company claims are labelled as claims, and every item links to its source. Spot an error or want a topic covered? Just reply. ✉