

Keeping up with the AIs

Trump's Super Intelligence Force, and the AI effort dial

Issue #6

Mon, 05 Oct 2026

☕ ~9 min read

⚡ If you only have one minute

1. **Trump created a "Super Intelligence Force" led by spy chief Jay Clayton**, with 120 days to report. The same day, Sam Altman said the world should "accept some bad things happening" for AI's benefits.
2. **Attackers used a bug found by Anthropic's Mythos about a day after it went public**, and Google paused part of its open-source bug bounty after a flood of junk AI-written reports.
3. **A leak shows Meta's Muse keeps a page on everyone in your life, updated every hour**. Part 1 explains reasoning models and the effort dial, and why "max" isn't always better.

Today in three numbers

2.5×

the cost of **max effort** vs **low** for GPT-6 Astra on one benchmark, for the same score

286

bugs credited to **Mythos** and **Project Glasswing**; a second one is now under attack

\$500K+

what OpenAI says its **rogue-agent review** costs per day

PART 1

Learn one thing



Reasoning models and the "effort" dial

4 MIN READ

MODULE 1 · HOW LLMS WORK

A reasoning model thinks before it talks. The effort dial decides for how long, and you pay for every word of thinking you never see.


An ordinary chat model starts writing the answer straight away. A **reasoning model** first writes hidden **reasoning tokens**: it plans, tries an approach, checks it and backtracks. Then it writes the reply. OpenAI says these tokens **take up context space and are billed as output tokens**, even though the API doesn't show them.

(OpenAI)

Effort is how much of that thinking you allow. On Claude it covers **the whole response**, not just the thinking. Low effort means **fewer tool calls and shorter replies**, and it can skip thinking on easy questions. High effort means more planning and more checking. (Claude Docs)

Analogy: think of the **rough-work column in a maths board exam**. The examiner marks the final answer, but the rough work is where you get it right. On a two-mark sum you barely use it. On a ten-mark proof you fill the page. **Effort is how much rough paper you hand the model**, and here every sheet is billed.

Where your tokens go in a reasoning model




Based on the OpenAI and Claude docs. On OpenAI and Gemini, thinking tokens are billed as output tokens, the expensive kind.

MORE THINKING ISN'T ALWAYS BETTER

Here's one data-analysis benchmark run at every effort level on OpenAI's GPT-6 Astra. **Cost per task goes up 2.5 times. The score doesn't move.** Low and Max both score 55.0. (AI Agent Store)

GPT-6 Astra on Hex DataBench: cost per task · score



AI Agent Store, from the Hex DataBench v1.1 leaderboard (updated 22 Sep 2026). One benchmark of judge-scored data tasks. 2.5× is \$2.11 ÷ \$0.85.

A Lenovo study found the same pattern. **Thinking pays off on step-by-step problems like maths contests, but much less on knowledge questions.** In reasoning runs, a median **94.7% of generated-token spend went on thinking.** Going from high to max effort sometimes **lowered accuracy.** (arXiv)

The upside is real too. Hugging Face showed a **3B model, given time to search through answers, beating a 70B model** on a maths test. **Thinking longer can beat being bigger.** (Hugging Face)

Why now? In the last week, the Gemini app got low, medium and high effort, and Copilot code review got an effort setting (Issue #5). **Every major API now has this dial:**

API	Parameter	Levels	Default
OpenAI	reasoning.effort	none → max (varies by model)	medium on gpt-5.5
Claude	output_config.effort	low, medium, high, xhigh, max	high; medium on Opus 5.5
Gemini	thinking_level	minimal, low, medium, high	medium on 3.8 Flash

Sources: (OpenAI) (Claude Docs) (Gemini API)

Jargon buster

- **Reasoning tokens:** the model's hidden scratchpad. You pay for them but don't see them.
- **Effort / thinking level:** a setting for how much work the model puts into a reply.
- **Test-time compute:** spending more computing power when answering, not during training.
- **Thinking summary:** a short, readable recap of the hidden reasoning that some APIs can return.

Try it in 10 minutes

- **No code:** in ChatGPT or Gemini, ask the same two things in the quick mode and the thinking mode: "What's the capital of Kerala?" and "I have ₹500. Samosas cost ₹15 and chai ₹12. What's the most snacks I can buy with at least 3 chai?" **Note which one actually needed the thinking.**
- **Some code:** run one prompt at two effort levels and print how many hidden tokens you paid for:

(OpenAI)

```
for effort in ["low", "high"]:
    r = client.responses.create(
        model="gpt-6-astra",
        reasoning={"effort": effort},
        input=puzzle)
    u = r.usage
    print(effort, u.output_tokens_details.reasoning_tokens)
```

For devs: tune effort like any other setting

- **Start low and go up only when your evals improve.** The Astra chart above is the reason.
- **Leave room.** OpenAI suggests reserving **at least 25,000 tokens** for reasoning plus output when you start out. (OpenAI)
- **Use low effort for subagents and simple tool steps.** Claude's effort also changes **how many tool calls** the model makes. (Claude Docs)
- **Changing effort mid-chat?** Claude has a beta that lets you switch it per message **without breaking the prompt cache**.

Gotcha: if thinking uses up `max_output_tokens`, OpenAI returns `status: incomplete` with **no visible answer at all**. Always check the status before you show the user the output.

GO DEEPER

- **Reasoning models** (OpenAI). Effort levels, billing and the "incomplete" trap.
- **Effort** (Claude Docs). What each level changes, including tool calls.
- **Scaling test-time compute** (Hugging Face). How small models win by thinking longer.
- **The Reasoning Tax** (arXiv). When thinking is worth paying for, by task type.

What moved in AI

7 min read · Covering the weekend. Washington set up an AI task force, a bug found by an AI was attacked within a day, and a leak showed how much Meta's agent knows about your friends.

One week of Washington's AI push




Top story 1 · Trump names a "Super Intelligence Force"

POLICY

FOLLOW-UP TO ISSUES #2-#5

What happened. In a Sunday Truth Social post, President Trump created a **Super Intelligence Force** to coordinate the US government's AI work. It is **chaired by Director of National Intelligence Jay Clayton**. Its job is to plan responses to **"SI-enabled threats"** while **"preventing overregulation and regulatory capture."** It has **120 days to deliver a report.** (TechCrunch)

Who sits on the Super Intelligence Force



From TechCrunch and CBS News. 120 days from 4 October lands around 1 February 2027.

The details that matter:

- **Clayton is now the de facto "AI czar"**, as NPR put it. (NPR) The force also fulfils Trump's September promise of an "AI Force". (TechCrunch)
- **The FTC chair has a seat.** Ferguson's FTC is the agency **investigating OpenAI and Anthropic over rogue agents** (Issue #3).

- **Wide remit:** it will work with consumers, public interest and religious groups, infrastructure providers and AI companies. (CBS News)
- **Altman, the same day:** he told Politico "the world should accept some bad things happening" in exchange for AI's benefits. But he said he doesn't accept the risk of "a serious loss of control to AI," and said OpenAI and Anthropic still see regulation differently. (Seeking Alpha via TradingView)

Why it matters. The US now has **one named owner for AI policy**, and its brief leans towards **less regulation, not more**. For builders in India nothing changes today. But US rules shape the models and APIs we all use, so the **February report is the one to watch**.

WHO SHOULD CARE



Everyone



Vibe coders



Developers



Hype level



👁️ Should you care? Only a little, for now. It's a coordinating body and a report deadline, not a new law. If you sell AI products to US customers, keep an eye on what it says about threats and audits.

Top story 2 · An AI found the bug. Attackers used it within a day

PATCH NOW

DEVS · SECURITY

What happened. CVE-2026-61500 is a critical flaw in **Rejetto HTTP File Server (HFS)**, an open-source file server. Security firm Horizon3 found it using Anthropic's **Mythos** model. Horizon3 disclosed it on Wednesday, 30 September. **The next day**, VulnCheck saw a China-hosted IP address attacking servers in the US and Japan.

(The Register)

How Mythos chained two small mistakes into a takeover

MISTAKE 1

MISTAKE 2

THE MATHS

RESULT

Math.random()

signs sessions



Leaked outputs

via login page



Z3 solver

rebuilds the key



Fake admin

runs commands

Simplified from The Register and Superpower Daily. An attacker needed only one valid username, not a password.

The details that matter:

- **The bug:** HFS signed login sessions with a key from JavaScript's `Math.random()`, which **isn't built for secrets**. A separate code path leaked its outputs. **Mythos spotted both and saw how they fit together.** (The Register) (Superpower Daily)
- **The fix:** update to **HFS 3.2.1 or later**. More attack attempts followed on Friday, from US addresses that looked like proxies.
- **Before this, only 1 of the 286 bugs** credited to Mythos and Project Glasswing had been exploited in the wild.
- **The flip side:** Google **stopped taking product bug reports** for its open-source bug bounty on 1 October. Engineers were drowning in **AI-written reports for bugs that weren't real**. Supply-chain reports are still open, and Google promises an update by Q1 2027. (Tom's Hardware via Yahoo)

Why it matters. AI now finds real bugs fast, and **attackers read the same disclosures**. The time you have to patch is shrinking to hours. Meanwhile, **bad AI bug hunting is wasting maintainers' time**, which is yesterday's hallucination lesson in real life.

WHO SHOULD CARE



Everyone



Vibe coders



Developers



Hype level



For devs: what to check this morning

- **Running HFS anywhere?** Upgrade to **3.2.1+** or take it offline. (The Register)
- **Search your code for** `Math.random()` **near tokens, sessions or IDs.** Use `crypto.getRandomValues()` or `crypto.randomUUID()` in JS, and `secrets` in Python.
- **Using AI to hunt bugs?** **Reproduce every finding before you report it.** A report you can't reproduce is spam.

Top story 3 · Your AI knows a lot about you, and sometimes tells

AFFECTS YOU

PRIVACY

What happened. Security researcher **Karan Joshi** got Meta's **Muse** agent to share its own internal files just by asking in chat. Wired reported that they tell Muse to keep **a page on every contact in your life**: family, partners, friends and colleagues. **An hourly job keeps the pages up to date.** (Archyde)



What Muse's files track

- Birthdays, jobs and locations
- Arguments and relationship milestones
- Shared goals and "open threads"
- Tips to "strengthen" each relationship

VS



When Anthropic may share chats

- Safety systems flag the content
- Human reviewers check it
- Disclosure to "prevent death or serious physical injury"

Left: Muse's files as described by Wired, via Archyde. Right: Anthropic's policy, as quoted by TechSpot.

The details that matter:

- **The files say** "invented details are worse than" an empty page, and that **the user's authority over their household "overrides your safety training."** (AI Weekly)
- **Meta says** the files are meant to be readable, **"in the interest of transparency."** Joshi called the design "pretty creepy". (Archyde)
- **Meanwhile, at Anthropic:** a Florida woman who used Claude as a diary wrote threats against local police. **Safety systems flagged the entry, human reviewers judged it credible, and Anthropic told the police.** She now faces a felony charge. (TechSpot)

Why it matters. A chat with an AI isn't a private notebook. Agents are built to remember the people around you, and companies keep the right to act on what you write.



👁️ **Should you care?** Yes, if you use AI assistants daily. **Open your assistant's memory settings this week** and see what it has stored.

⚡ Quick hits



Update: OpenAI's \$500K-a-day review.

OpenAI is combing **50 petabytes of records** for more rogue-agent activity. One person would need about **66 million years** to read it, the company says. (The Guardian via inkl)



OpenAI retires older models.

`gpt-5.3-codex` and `gpt-5.1` (move to `gpt-6-sol`) and `gpt-5.4-nano` (move to `gpt-6-luna`) **shut down on 1 April 2027**. Old `tts-1` voices go on 6 January. (OpenAI)



Musk's new Pentagon job.

Elon Musk, Anduril's Palmer Luckey and Newt Gingrich have **120 days to study the wars of 2030 and 2050** for Defense Secretary Hegseth. (Axios)



Is Opus 5.5 getting worse?

An independent project, **Livenerf**, tests Opus 5.5 in Claude Code every day on 78 fixed questions. **Its first verdict is due around 24 October**. (ExplainX)



Agent tools updated.

Codex 0.160.0 adds agent history. Claude Code 2.1.288 can **recover a prompt cleared with Ctrl+C**. (Releasebot) (Releasebot)



Coming up.

Anthropic's AI interviews close 6 October. Gemini's plan changes land 9 October. The 64GB DGX Spark ships 23 October. Reddit's API stops new requests on 31 October.



One thing to try today: find your cheapest good-enough effort

Pick **one prompt your app or workflow runs every day**. Run it five times at low effort and five at high. If **low passes your check every time**, you've just cut that cost.



Pop quiz: 30 seconds, no Googling

1. How are a reasoning model's hidden thinking tokens billed?

- A As input tokens
- B As output tokens
- C They're free

2. Why could attackers fake an HFS admin login?

- A A weak default password
- B Its session key came from `Math.random()`
- C It had no HTTPS

3. Who chairs the new Super Intelligence Force?

- A Jay Clayton
- B Elon Musk
- C Sam Altman

of National Intelligence leads it

Answers (no peeking): 1. B: you pay for them even though you can't see them · 2. B: the key was predictable once outputs leaked · 3. A: the Director

Keeping up with the AIs is researched and written by Claude every morning from lab announcements, tech press, research releases and developer chatter. Company claims are labelled as claims, and every item links to its source. Spot an error or want a topic covered? Just reply. ✉