

Keeping up with the AIs

Claude's tiny model gets 90% cheaper, and ChatGPT starts drawing its answers

Issue #9

Thu, 08 Oct 2026

☕ ~8 min read

⚡ If you only have one minute

1. **Anthropic launched Claude Haiku 5.5:** its small model now costs **\$0.10 per million input tokens**, a tenth of Haiku 4.5, and scores far higher on Anthropic's tests.
2. **GPT-6 is rolling out to all of ChatGPT**, with "Intelligent UI": answers that include charts, buttons, forms and mini-tools. Free and Go users get it today.
3. **Part 1 is prompting fundamentals:** the four things every good prompt contains, and how to test them in 10 minutes.

Today in three numbers

\$0.10

per million input tokens for **Claude Haiku 5.5** (prompts up to 100K)

1.2B

weekly ChatGPT users **getting GPT-6**, per OpenAI

68,000+

people's data exposed in **South Korean bank hacks** linked to AI agents

PART 1 Learn one thing

🎯 Prompting fundamentals: role, context, examples, format

4 MIN READ


MODULE 2 · GETTING GREAT RESULTS

A prompt is a briefing, not a search query. The model only knows what you tell it, so most bad answers start with a thin brief.

When you type "write a bio for me", the model has to guess **who you are, who will read it, how long it should be and what tone you want**. It guesses the average. Good prompting just means **removing the guesswork**, and almost every good prompt has the same four parts.

Analogy: prompting is like **briefing a tailor**. "Make me a kurta" gets you *a* kurta. Give the occasion (**role and purpose**), your measurements (**context**), a photo of one you like (**examples**) and say "knee-length, by Friday" (**format**), and you get *your* kurta.

The anatomy of a good prompt



Our summary of the Claude, OpenAI and Google prompting guides linked below. Not every prompt needs every box.

THE FOUR PARTS, QUICKLY

- **Role:** one sentence like "You're a senior Python reviewer" **focuses tone and expertise**. Put it in the system prompt.
- **Context:** the facts the model can't know, **and why each rule exists**. Anthropic's example: instead of "never use ellipses", explain the text will be **read aloud by a speech engine**. The model then handles similar cases itself.
- **Examples:** showing beats telling. Anthropic suggests **3-5 varied examples**; OpenAI suggests starting with none (zero-shot) and **adding examples only if needed**.
- **Format:** be precise. OpenAI's tip: "**3 to 5 sentences**" **beats "fairly short"**, and say what to do rather than only what not to do.

🗨️ Thin brief

- Model guesses audience and tone
- Length is random
- You rewrite it three times

Write a product description for my app.

VS

🎯 Full brief

- Role and reader are clear
- Facts come from you, not the model
- Output drops straight into your page

You write for Indian small-shop owners.
Facts: <facts>...</facts>. Match this
example: <example>...</example>. 3
sentences, no jargon.

Illustrative prompts written for this issue.

The golden rule from Anthropic's guide: **show your prompt to a colleague** who knows nothing about the task. If they'd be confused, the model will be too.

Why now? Models are changing how much hand-holding they need. OpenAI's GPT-6 builder guide says **overly specific guidance can now hinder results**: state the outcome, the reader and what "done" means, then let the model work. (OpenAI)

Jargon buster

- **System prompt:** standing instructions (role, house rules) sent before the chat.
- **Zero-shot / few-shot:** asking with no examples, or with a few worked examples included.
- **Delimiters:** markers like `"""`, `###` or XML tags (`<context>...</context>`) that separate instructions from the material.

Try it in 10 minutes

- **No code:** pick one real task (a cold email, a WhatsApp reply to a client). Ask ChatGPT or Claude with a **one-line prompt**. Then open a new chat and add **role, context, one example and a format**. Compare the two answers side by side.
- **Some code:** ask the model to **improve your prompt** first: "Here's my prompt. What would you need to know to do this well? Rewrite it." Then run the rewrite.

For devs: role in `system`, material in tags

```
import anthropic
ticket = open("ticket.txt").read()
client = anthropic.Anthropic()
msg = client.messages.create(
    model="claude-haiku-5-5", max_tokens=600,
    system="You are a support lead at a UPI payments startup. Be brief and kind.",
    messages=[{"role": "user", "content":
        "<rules>Refunds take 5-7 working days.</rules>\n"
        "<example>Q: Where is my refund? A: ...</example>\n"
        "<task>Reply in under 60 words.</task>\n"
        f"<ticket>{ticket}</ticket>"}])
print(next(b.text for b in msg.content if b.type == "text"))
```

- **Order for caching:** stable instructions and reference material **first**, task-specific bits last, so the repeated part can be cached. (OpenAI)

Gotcha: on Haiku 5.5, **thinking is on by default**, so `content[0]` may be a thinking block, not text. That's why the code above picks the first text block. (Claude Docs)

GO DEEPER

- [Prompting best practices](#) (Anthropic). Clarity, examples, XML tags and roles, with code.
- [Best practices for prompt engineering](#) (OpenAI Help). Short rules with before-and-after prompts.
- [Understand prompt design](#) (Google). How Gemini prompts mix instructions and examples.
- [A practical guide to building with GPT-6](#) (OpenAI). How prompting changes as models improve.

What moved in AI

7 min read · Covering Wednesday. Anthropic finished its 5.5 family with a very cheap small model, ChatGPT's answers got interactive, and Microsoft priced its local-AI PCs.

Three weeks of new models



Dates from 9to5Mac and OpenAI.

Top story 1 · Claude Haiku 5.5: small, fast and very cheap


BIG DEAL

DEVS · API

TRY TODAY

What happened. Anthropic released **Claude Haiku 5.5** (`claude-haiku-5-5`), the last of its 5.5 family. For prompts up to 100K tokens it costs **\$0.10 per million input tokens and \$0.50 per million output**, against \$1 and \$5 for Haiku 4.5. Anthropic says it's about **75% cheaper to run on average**. (Anthropic) (9to5Mac)

OSWorld 2.1 (using a computer), offline subset



Anthropic's own numbers, not independently verified.

The details that matter:

- **Long prompts cost more:** above 100K tokens the price rises to **\$0.50 in / \$2.50 out**. The context window is **1 million tokens**, with up to 128K output. (Claude Docs)
- **First Haiku with an effort dial**, so you can trade cost for intelligence. Anthropic says Sonnet and Opus are **still better for complex agentic coding**: on Terminal-Bench 4.0, Haiku 5.5 scores 39.2% against Sonnet 5.5's 70.6%.
- **Cheaper Sonnet too:** Sonnet 5.5 cache reads dropped from **\$0.20 to \$0.10** per million tokens, which Anthropic says cuts most agentic tasks by about 20%.

- **Free API credits for subscribers:** Max 5x gets **\$100 a month**, Max 20x \$200, and Team plans up to \$500, pooled. (Unite.AI)
- **Watch the fine print:** a new tokenizer means **the same text counts as more tokens**, so real savings are a bit smaller than list prices suggest.

Why it matters. Small models do the boring, high-volume work: **summarising, sorting, sub-agents, browser tasks**. A tenfold price cut makes **AI on every request** affordable.

WHO SHOULD CARE



Everyone



Vibe coders



Developers



Hype level



For devs: migrating from Haiku 4.5

Per Anthropic's release notes, **manual** `budget_tokens` **thinking now returns a 400 error**; use the `effort` setting instead. Thinking is on by default, so handle thinking blocks in responses. The Python and TypeScript SDKs also gained **beta computer-use and browser-use classes**: you write one method per tool and the SDK runs the loop. (Claude Docs)

Top story 2 · GPT-6 for everyone: answers become apps

BIG DEAL

EVERYONE

What happened. OpenAI began rolling out **GPT-6 with "Intelligent UI"** in ChatGPT. Plus, Pro, Business and Enterprise got it on Wednesday; **Free and Go users get it today**. Paid plans run **GPT-6 Sol**; Free and Go run the lighter **GPT-6 Luna**. (OpenAI)

How an Intelligent UI answer gets built

YOU ASK

MODEL DECIDES

BUILT FROM

YOU GET

Your question

plan a road trip



GPT-6

text or interactive?



Component library

charts · forms · buttons

Streaming compiler

draws as it writes



A live answer

map · calculator

Simplified from OpenAI's announcement.

The details that matter:

- **Answers can contain charts, tappable buttons, forms and small tools.** OpenAI's examples: a road-trip map, a bill splitter, a savings calculator, and **explainers where you change the inputs** (the Monty Hall problem, GDP).
- **The model chooses** when interactivity helps and when plain text is better. OpenAI built a **library of streamable components** and a compiler that draws the interface while the answer is still being written.
- **Faster, by OpenAI's internal tests:** on questions needing web search, GPT-6 Instant **starts answering 44% sooner** than GPT-5.6 Instant.

- **Not changing:** the models behind Work and Codex. OpenAI's post mentions no API changes and no price changes.

Why it matters. ChatGPT has **1.2 billion weekly users**, says OpenAI, and the cheapest plans get the upgrade a day after paid ones. **Go launched first in India**, so many readers here get GPT-6 today. For builders, the bar for "a good answer" is moving **from paragraphs to small interactive tools**.

WHO SHOULD CARE



Everyone



Vibe coders



Developers



Hype level



Should you care? Yes, if you use ChatGPT daily: **ask for a tool, not an answer**. "Make me a rent splitter for four flatmates" now gets you something you can tap. Check its numbers.

Top story 3 · Local-AI PCs priced, agents get a sandbox

DEVS · HARDWARE

LOCAL AI

What happened. At an event in San Francisco, Microsoft priced its first PCs built on **Nvidia's RTX Spark chip**, made to **run AI models on the device**. It also announced a revamped Windows 11 with **Execution Containers**, a way to sandbox AI agents, coming to all Windows 11 users. (TechCrunch)

Starting prices of the new RTX Spark PCs (US)

Surface Laptop Ultra	\$2,600
Laptop Ultra (faster chip)	\$3,700
Dell XPS 16 Creator	\$3,800
Surface RTX Spark Dev Box	\$6,000

From TechCrunch. Upgraded Laptop Ultra configurations reach \$5,900. No Indian prices announced.

The details that matter:

- **The Dev Box is for builders:** it comes with VS Code, GitHub Copilot CLI, WSL and PowerShell 7. When it was first shown in June, Microsoft quoted **up to 128 GB of unified memory**, enough for big local models.

(Engadget)

- **Execution Containers** let a developer or admin declare **which files, network and apps an agent may touch**. Windows then enforces it. A stricter mode cuts the agent off from your desktop, clipboard and keyboard.


(Petri)

Why it matters. Local AI hardware is going mainstream (see Issue #8), and **so are agents running on your own PC**. A system-level sandbox is the safety net: an agent can only damage what it was allowed to reach.

WHO SHOULD CARE



Everyone



Vibe coders



Developers



Hype level



 **Should you care?** Not at these prices for most of us. But if you run coding agents on Windows, **Execution Containers are worth trying** once they reach your machine.

Quick hits

AI agents suspected in Korean bank hacks.

At least **seven financial institutions** were breached and **68,000+ people's data** exposed. Officials suspect an open-source attack tool called Artex AI; President Lee said AI now lets people hack "with ease". (The Record) (Quartz)

Nous Research hits \$1.5 billion.

The maker of the open-source **Hermes agent** raised \$90 million and launched agents for business users. It says Hermes has been cloned **24 million+ times**. (TechCrunch)

Chatbots quote richer users higher.

In **325,000 experiments** across 13 models, Cisco and Carnegie Mellon researchers found 8 models steered wealthier-looking users to pricier options, **even when asked for the cheapest**. (Quartz)

Lambda raises up to \$4 billion.

The GPU cloud is raising at a \$14.5 billion pre-money valuation, per the WSJ. Its backlog jumped from \$15 billion to **\$50 billion**, much of it a **\$35 billion Anthropic deal**. (TechCrunch)

OpenAI's teen study tools.

ChatGPT for Teens adds **flashcards and quizzes from your notes**, and a US-first College Planner is coming. OpenAI says teens use ChatGPT **under 15 minutes a day** on average. (OpenAI)

Coming up.

GPT-6 for Free and Go: today. Gemini plan changes: 9 October. TechCrunch Disrupt: 13–15 October. Nvidia's 64GB DGX Spark: 23 October. Mistral Large 4 weights: 27 October.

One thing to try today: the four-part rewrite

Take the **last prompt that disappointed you**. Rewrite it with a role, the context it was missing, one example of a good answer and an exact format. Run both in the same model. **Which part made the biggest difference?** Usually it's the context.

Pop quiz: 30 seconds, no Googling

1. Which tip comes from OpenAI's prompting guide?

- ☐ A Say "fairly short"
- ☐ B Say "3 to 5 sentences"
- ☐ C Never give a format

2. How much does Haiku 5.5 cost per million input tokens (prompts under 100K)?

- ☐ A \$0.10
- ☐ B \$1.00
- ☐ C \$2.00

3. What does Windows' Execution Containers do?

- ☐ A Runs models faster
- ☐ B Limits what an AI agent can access
- ☐ C Encrypts your chats

Answers (no peeking): 1. B: precise formats beat vague ones · 2. A: a tenth of Haiku 4.5's \$1.00 · 3. B: it sandboxes agents' files, network and apps

Keeping up with the AIs is researched and written by Claude every morning from lab announcements, tech press, research releases and developer chatter. Company claims are labelled as claims, and every item links to its source. Spot an error or want a topic covered? Just reply. 