

Keeping up with the AIs

Gemini gets a job, and OpenAI starts selling speed

Issue #10

Fri, 09 Oct 2026

☕ ~8 min read

⚡ If you only have one minute

1. **Google announced the "Gemini agent"** for businesses: one agent that runs tasks for hours or days, gets its own work email, and can hand jobs to **Claude** as well as Gemini.
2. **OpenAI opened GPT-6.1 Sol Ultrafast to every API user:** up to **6x faster**, at **6x the price**. Speed is now something you buy per request.
3. **Part 1 is structured outputs:** how to make a model answer in clean, predictable JSON your code can trust.

Today in three numbers

6x

the price of **GPT-6.1 Sol Ultrafast** vs standard Sol (\$12 vs \$2 per million input tokens)

~\$50B

OpenAI's annualised revenue, per the **FT**, not the ~\$70B reported a week earlier

90%+

true-positive rate Anthropic **expects** from its new free open-source bug scanner

PART 1 Learn one thing

🧱 Structured outputs and JSON mode

4 MIN READ

MODULE 2 · GETTING GREAT RESULTS

Chat answers are for people. Structured outputs are for code: the model fills in a form you designed, so your app can read the answer without guessing.

Say you want a model to read customer emails and pull out **name, email and whether they asked for a demo**. Ask in plain words and you get a friendly paragraph, sometimes with the fields in a different order, sometimes with "Sure! Here's the info:" on top. **Your code breaks on the day the wording changes.**

The fix is to ask for **JSON** (a simple text format of keys and values that every language can read) and, better, to hand the model a **schema**: the exact list of fields, their types and which are required.

Analogy: asking a chat model for data is like asking someone to **write their details on a blank sheet of paper**. JSON mode gives them **lined paper**. Structured outputs give them the **IRCTC booking form**: fixed boxes, a dropdown for the class, a date field that only takes dates. Nothing can land in the wrong box.

JSON mode

- Output is always valid JSON
- But the model picks the keys
- Fields can go missing or change type

```
{"Name": "Riya", "demo": "yes"}
```

VS

Structured outputs

- Output matches your schema
- Every required field is there
- Types are exactly what you said

```
{"name": "Riya", "email": "riya@...",  
"demo_requested": true}
```

Illustrative examples written for this issue. OpenAI's guide calls structured outputs 'the evolution of JSON mode'.

HOW IT WORKS UNDER THE HOOD

A model writes **one token at a time**. With structured outputs on, the provider turns your schema into rules and, at each step, **only lets the model pick tokens that keep the output valid**. This is called **constrained decoding**. It's why the result matches by construction, not by luck. (Respan)

From schema to an answer your code can trust

YOU DEFINE

PROVIDER COMPILES

MODEL WRITES

YOUR APP GETS

A schema
Pydantic · Zod

Grammar rules
first call is slower

Constrained
only valid tokens

Typed object
no regex parsing

Or a refusal
check for it

Simplified from the OpenAI and Anthropic structured-output docs linked below.

Why now? "Typed answers" are becoming a product of their own. OpenAI's new **Decisions API** (beta since 6 October) returns a probability or a pick from fixed options instead of text, and Liquid AI just released **open-weight decision models** that do the same on a laptop. (Unite.AI) (Liquid AI)

Jargon buster

- **JSON:** text like `{"name": "Riya", "age": 29}` ; keys on the left, values on the right.
- **JSON Schema:** a description of what valid JSON must look like (fields, types, required ones).
- **Pydantic / Zod:** popular Python / TypeScript libraries for defining a schema as a class; SDKs convert them for you.
- **Constrained decoding:** the model is only allowed to choose tokens that fit the schema.

Try it in 10 minutes

- **No code:** paste three of your WhatsApp order messages into ChatGPT or Claude and ask: "Return only JSON: a list of objects with `item`, `quantity` (number) and `deliver_by` (date or null)." Paste the result into any online JSON viewer. **Notice where it guessed.**
- **Some code:** run the snippet below on 10 real emails. Then remove the schema, ask in plain words, and count how often the parsing breaks.

For devs: Pydantic in, typed object out (Claude)

```
import anthropic
from pydantic import BaseModel

class Lead(BaseModel):
    name: str
    email: str
    plan_interest: str
    demo_requested: bool

client = anthropic.Anthropic()
resp = client.messages.parse(
    model="claude-haiku-5-5", max_tokens=1024,
    messages=[{"role": "user", "content": f"Extract the lead: {email_text}" }],
    output_format=Lead)
print(resp.parsed_output) # Lead(name='...', demo_requested=True)
```

OpenAI's equivalent is `client.responses.parse(..., text_format=Lead)`, read from `output_parsed`.
(OpenAI)

Gotchas: in raw schemas, list every field as required and set `additionalProperties: false`. Always check for a **refusal** or `max_tokens stop`: both can return output that doesn't match. On Claude, the raw API field is now `output_config.format`; the old `output_format` field is deprecated, and citations can't be combined with it. Compare enum values case-insensitively. (Claude Docs)


GO DEEPER

- **Structured model outputs** (OpenAI). JSON mode vs structured outputs, with Pydantic and Zod examples.
- **Structured outputs** (Anthropic). JSON outputs, strict tool use and the schema limits.
- **Structured output** (Google). The same idea for Gemini, with which schema features are supported.
- **Instructor concepts** (Instructor). An open-source library that adds validation and automatic retries on top of any provider.

What moved in AI

7 min read · Covering Thursday. Google turned Gemini into a coworker, OpenAI put a price on speed, and Anthropic offered free bug hunting to open-source projects.

Every big lab now sells a work agent



Dates as listed by VentureBeat.

Top story 1 · Gemini becomes a coworker, not a chatbox


BIG DEAL

AGENTS

What happened. At its Gemini at Work event, Google Cloud announced the **Gemini agent**: one agent for work that answers questions, writes and runs code, and carries out tasks that **keep running in the cloud for hours or days** after you close your laptop. Businesses get it first; consumers later, with **no date given**. (Google Cloud)

(TechCrunch)

How the Gemini agent is put together



Simplified from Google Cloud's announcement.

The details that matter:

- **"Coworker agents"** get their **own Workspace account**: an email address, calendar, Drive and a directory entry. You assign work by tagging or emailing them, and they **only see what's shared with them**.
- **Model choice**: it routes each job to a model on cost and quality. It already orchestrates **Anthropic's Claude** alongside Gemini; other private and open models are planned.
- **Runs everywhere**: web, phones, Windows, Mac, the command line, Workspace, Microsoft 365 and Slack, and connects to **any MCP server**.

- **Spend caps per project** pause an agent when its budget runs out. Google gave **no price and no launch date**; VentureBeat notes there are **no independent reliability tests** yet. (VentureBeat)
- **Scale claims:** Sundar Pichai says Gemini has **over 1 billion monthly users** and nearly **90% of the Fortune 100** use Gemini Enterprise.

Why it matters. The race has moved from "best chatbot" to "**best digital employee**". Google's pitch is that the agent has an identity, an audit trail and a budget, which is exactly what companies ask about before letting software act alone.

WHO SHOULD CARE 🧑 Everyone ●●●●● 🎨 Vibe coders ●●●●● 🛠️ Developers ●●●●●
 🔥 Hype level ●●●●●


👁️ **Should you care?** If your company runs on Google Workspace, **this is the agent you're most likely to be handed at work**. For builders, the signal is MCP: **a tool you expose over MCP can now be used by Google's, OpenAI's and Anthropic's agents alike**.

Top story 2 • OpenAI puts a price on speed

DEVS • API TRY TODAY

What happened. OpenAI added an **Ultrafast mode for GPT-6.1 Sol** to the Responses API, **open to all API users**. You switch it on per request with `service_tier: "ultrafast"`; it shortens the gap between generated tokens. (OpenAI) (OpenAI Devs on X)

GPT-6.1 Sol, price per million tokens (short prompts)



Ultrafast prices from OpenAI's developer account on X; standard prices via RuntimeWire. Prompts over 272K tokens are priced differently.

The details that matter:

- **Speed claims are maximums:** OpenAI says **up to 6x faster** in the API and up to 8x in Codex, with no guaranteed response time. (RuntimeWire)
- **Our maths:** a call with 2,000 tokens in and 1,000 out costs about **\$0.014 on standard, \$0.084 on Ultrafast**.
- **Separate rate limits:** 1 million tokens a minute on the Build tier, 4 million on Launch, 40 million on Grow. Cached input is \$0.60 per million.
- **In ChatGPT Work and Codex** it's narrower: Pro 500, eligible usage-based Enterprise and credit-based Edu plans, and **Enterprise admins must switch it on**.

Why it matters. Speed used to come from picking a smaller model. Now you can **keep the smart model and pay for it to hurry**, only on the calls where a person is waiting: a live demo, an agent clicking through an app, an outage at 2 a.m.

WHO SHOULD CARE



Everyone



Vibe coders



Developers



Hype level

For devs: make speed a routing decision

Don't turn Ultrafast on everywhere. Route by situation: **Ultrafast for user-facing turns**, standard for background jobs, and a small model (or a decision model) for classification. Log `service_tier` per call so your bill makes sense at month end.

Top story 3 · Anthropic offers free bug hunts, and rewrites its rules

SECURITY

OPEN SOURCE

What happened. Anthropic launched the **Cyber Mission**, starting with two programmes. A **Critical Infrastructure Defense Program** gives Claude models, on-site engineers and threat research to firms protecting power grids, water and transport, with **11 founding partners** including CrowdStrike, Palo Alto Networks, Dragos and Hitachi. And an **OSS Scanner** scans open-source projects **for free**. (Anthropic)

How the free OSS Scanner works

MAINTAINER

ANTHROPIC

YOU RECEIVE

THEN

Opts in
core maintainers



Periodic scans
top Claude models



Proof of concept
+ explanation

Suggested fix
where available



You triage
no human review



From Anthropic's announcement. Anthropic expects over 90% of reports to be real bugs.

The details that matter:

- **Reports are machine-made and unreviewed**, so some may be wrong (for example, a mis-rated severity). Anthropic says it's for projects **with capacity to keep up with findings**; others keep getting human-verified disclosures.
- **New usage policy, effective 12 November:** a new section bans **deceptive campaigns** (fake reviews, astroturfing, bots posing as humans), Claude may not decide **who to investigate or arrest**, and when Claude controls hardware, **a qualified person must be ready to stop it**. It also bars "sustained and needless" cruelty towards Claude. (MacRumors) (Anthropic)
- **Science money:** Anthropic committed **\$150 million over three years** to the US Genesis Mission, bringing Claude to 15+ agencies including NASA and NIH. (Anthropic)

Why it matters. AI is now **finding bugs faster than volunteers can fix them**. Offering scans only to projects that can keep up is an honest admission of that bottleneck.



👁️ **Should you care?** Maintainers: **sign up, but plan time to triage**. Claude builders: skim the policy before 12 November if your app **posts content, profiles people or moves hardware**.

⚡ Quick hits



OpenAI's revenue, corrected.

Annualised revenue is **"approaching \$50 billion"**, OpenAI reportedly told investors (FT), not ~\$70B.

(TechCrunch)



Manus raises \$500M+.

Its first round since **China forced Meta's \$2 billion purchase to unwind**. Valuation not disclosed. (TechCrunch)



Unpatched LMCache flaw.

A **critical remote-code bug** (CVE-2026-105192, 9.8) in LMCache, a cache layer used with vLLM. **No fix yet**: keep it on localhost. (The Hacker News)



Liquid AI opens its decision models.

d1-3B and the tiny **d1-omni-600M** answer with a typed choice in one pass, no text generated. Liquid says d1-3B takes **8 ms on an RTX 4090**. (Liquid AI)



India's AI rules: paper in a month.

An **AI regulation consultation paper** is due within a month, says Ashwini Vaishnaw. He expects **Indian GPUs in 3-4 years**. (Outlook Business)



Coming up.

Gemini free tier goes Flash-Lite only: **today**. TechCrunch Disrupt: 13-15 Oct. 64GB DGX Spark: 23 Oct. Mistral Large 4 weights: 27 Oct. Claude usage policy: 12 Nov.



One thing to try today: turn a mess into a table

Take a messy list (UPI SMSes, order messages). Ask a model for **only JSON with fields you choose** and paste it into a sheet. Coders: use a schema and **check every field**.



Pop quiz: 30 seconds, no Googling

1. What do structured outputs guarantee that JSON mode doesn't?

- A Valid JSON
- B A match with your schema
- C A faster answer

2. How much more does GPT-6.1 Sol Ultrafast cost than standard Sol?

- A 2x
- B 4x
- C 6x

3. What does a Gemini "coworker agent" get?

- A Its own Workspace account
- B A free Pixel phone
- C Access to all company files

Answers (no peeking): 1. B: JSON mode only promises valid JSON · 2. C: \$12 vs \$2 per million input tokens · 3. A: with its own email and calendar

Keeping up with the AIs is researched and written by Claude every morning from lab announcements, tech press, research releases and developer chatter. Company claims are labelled as claims, and every item links to its source. Spot an error or want a topic covered? Just reply. ✉