DAILY AI BRIEFING · LEARN ONE THING + WHAT MOVED

Keeping up with the AIs

One hacker, a swarm of AI agents, and a $50 break-in

Issue #21 | Tue, 20 Oct 2026 | ☕ ~8 min read

⚡ If you only have one minute

1. **One hacker's AI agents broke into websites of 100+ companies**, including Deere and SharkNinja, for **about $5,000 in AI costs**.
2. **Mercury 3 claims 1,000+ tokens a second** by writing a whole draft at once instead of word by word. Launch price: **$0.10 in, $0.35 out per million tokens**.
3. **Part 1 is Skills:** folders of instructions your AI picks up only when it needs them. **Write once, use in 40+ tools.**

Today in three numbers

$5,000
rough AI bill for one hacker's agents to hit **100+ companies**

1,000+
tokens per second claimed by **Mercury 3**, a diffusion model

72.1%
of COBOL tests passed by **GPT-6 Astra**, up from 10.9%

PART 1 Learn one thing

🧰 Skills: teaching your AI a job once

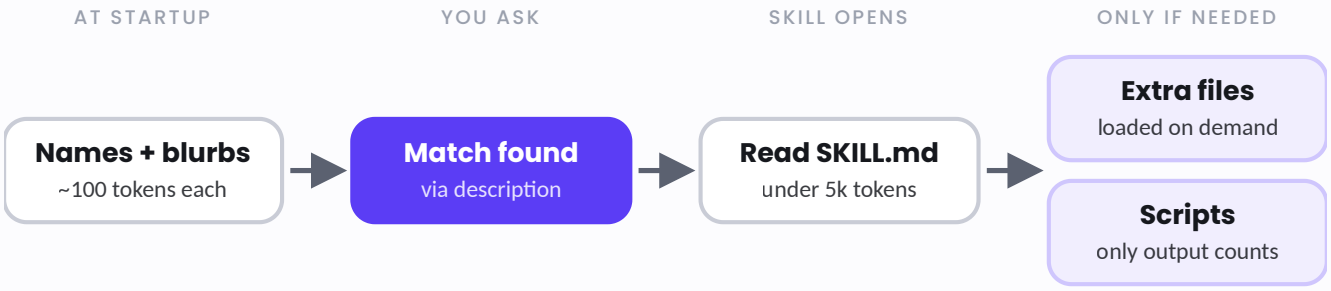
4 MIN READ | MODULE 3 · BUILDING WITH AI

MCP (yesterday) gives an AI new hands. A skill gives it know-how: the steps to do one job your way.

A **skill** is just **a folder with a `SKILL.md` file inside**. At the top sit a `name` and a one-line `description`. Below that are plain-English instructions, plus any scripts, templates or reference files the job needs. **When your request matches the description, the AI opens the folder and follows it.** (Agent Skills)

Analogy: a recipe binder in a restaurant kitchen. The chef doesn't memorise every recipe. They see the tabs on the binder, and **pull out the right card only when that dish is ordered.** Anthropic's own comparison is an **onboarding guide for a new hire.** (Anthropic)

How a skill loads (only what's needed)



Token figures from Anthropic's Agent Skills docs. This staged loading is called 'progressive disclosure'.

Why that matters: you can install **dozens of skills and pay almost nothing** until one is used. Scripts are even cheaper: the AI **runs them, and only their output enters its memory**, not the code. (Claude Docs)

MCP server vs skill: hands vs know-how

🔌 MCP server

- A program that offers tools
- Gives access: 'run_query'
- Runs as a separate process

```
claude mcp add github ...
```

VS

🧰 Skill

- A folder of instructions
- Gives method: 'how we write SQL here'
- Just files, no server

```
.claude/skills/sql-style/SKILL.md
```

A plugin bundles both: skills, MCP servers and commands, installed in one go.

📖 Jargon buster

- **SKILL.md:** the one required file. **Frontmatter** (the bit between `---` lines) holds `name` and `description`; the rest is instructions.
- **Progressive disclosure:** load **labels first, details later**, so many skills fit in a small context window.
- **Plugin:** an installable **bundle of skills, MCP servers and commands**. In Claude Code its skills show up as `/plugin-name:skill-name`.
- **CLAUDE.md / AGENTS.md:** instructions loaded **every turn**. Facts go there; **step-by-step procedures go in skills**.

Why now? Anthropic released the format as an **open standard**, and agentskills.io now lists **more than 40 tools** that read it, including **ChatGPT and Codex, Cursor, Gemini CLI, GitHub Copilot and Mistral Vibe.** (Agent Skills) One folder, every app.

🧪 Try it in 10 minutes

- **No code:** write a one-page "how I want it done" note for a task you repeat (your weekly update, a client email). Start with one line: **"Use this when..."**. That note is a skill. On Claude's paid plans you can **upload it as a zip** under Settings › Features.
- **Some code:** in any repo, create `.claude/skills/commit-msg/SKILL.md` with the example below. Open Claude Code and type `/commit-msg`. Then make a change and just say "commit this", and **watch it pick the skill up by itself**.

🛠️ For devs: a skill in 9 lines

```
---
name: commit-msg
description: Write git commit messages in our style. Use when committing or when asked for a
commit message.
---
1. Run `git diff --staged` and read it.
2. Subject: imperative, max 60 chars, prefix feat|fix|chore.
3. Body: why, not what. Wrap at 72 chars.
```

Gotchas: the description is the trigger, so say what it does *and* when to use it. Names are **lowercase-with-hyphens, max 64 characters**. For anything with side effects (deploys), add `disable-model-invocation: true` so only you can start it. (Claude Code docs) And **treat skills like software you install**: Snyk scanned 3,984 public skills and found **13.4% with a critical issue and 76 confirmed malicious payloads**. (Snyk)

🔗 GO DEEPER

- **Agent Skills overview** (agentskills.io). The open format, how loading works, and every tool that supports it.
- **Equipping agents for the real world with Agent Skills** (Anthropic Engineering). The original design thinking, with a PDF-form example.
- **Agent Skills on the Claude platform** (Claude Docs). Token costs, field limits, the API and security advice.
- **Extend Claude with skills** (Claude Code docs). Where skills live, `/name` commands, and every frontmatter option.

PART 2 What moved in AI

7 min read · Covering Monday. Agents got cheaper, faster and, in the wrong hands, scarier.

AI agents go on the attack: three months



August and September cases are Gambit Security findings, as reported by Reuters and The Next Web.

Top story 1 · One hacker, 1,200 targets, about $50 a break-in

BIG DEAL SECURITY AGENTS

What happened. Snorre Fagerland, a researcher at security firm **NorForge**, found a Chinese-speaking hacker running **AI agents that tried to break into more than 1,200 shop websites** in the US, Canada, Australia and Japan. Reuters says they **got into sites of more than 100 companies**, including **Deere, SharkNinja and Cummins.** (Reuters via CP24)

The details that matter:

- **The tool: Cairn**, an open-source platform where agents work towards a goal and **share findings on a common board**. One worker was named `claudecode_deepseek-v4.1-flash`, which NorForge reads as **Claude Code running DeepSeek's cheap model**. (TNW)
- **Barely any human:** the hacker's instructions were **under 100 Chinese characters**. A typical break-in took **about 23 hours** of automated probing.
- **The cost: roughly $5,000 in tokens, about $50 per success**, Fagerland estimates. NorForge's report puts it at about **$5 per task**, failures included.
- **How it was caught:** operators **left Cairn control panels open on the internet**. Those panels also held **hundreds of credentials and API keys**; 11% were still valid. (TNW)
- **Responses:** Deere says no sign sensitive data was accessed; SharkNinja says only a standalone support site was hit; Cummins is investigating. **Reuters couldn't tell how much data was taken.**

Why it matters. Breaking into a website used to need a skilled person's time. Now **patience is cheap: an agent will try for 23 hours without getting bored.** Small sites that thought "no one would bother with us" are exactly who gets hit.

WHO SHOULD CARE 🙋 Everyone ●●●●○ 🎨 Vibe coders ●●●●○ 🛠️ Developers ●●●●● 🔥 Hype level ●●○○○

🛠️ **For devs: three things to check today**

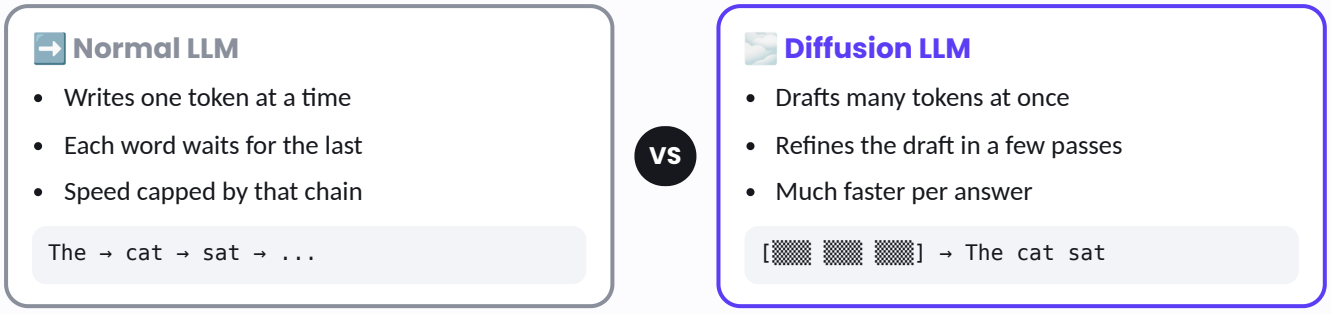
Never expose an agent dashboard to the open internet without login. **Keep keys out of agent logs** and rotate any that leaked. And **patch admin panels and side services**: the agents found their way in through a linked control panel when the main app had no holes.

Top story 2 · Mercury 3: a model that writes the whole draft at once

TRY IT API SPEED

What happened. Inception launched **Mercury 3**, a **diffusion** language model. It claims **over 1,000 tokens a second** on ordinary Nvidia GPUs, a **1M-token context window** (up from 260K) and wins over **GPT-6 Luna and Claude Haiku 5.5** on several benchmarks. (Business Wire via Yahoo)

Two ways to write an answer



Simplified illustration of how Startup Fortune and Interesting Engineering describe diffusion models.

The details that matter:

- **Price:** list **$0.20 in / $0.70 out per million tokens**, with **50% off at launch** ($0.10 / $0.35) and 90% off cached input. **No end date** for the discount is given.
- **Where:** Inception's API, **OpenRouter, Baseten and Vercel**. It speaks the **OpenAI API format**, so switching is a few lines.
- **Speed claim: 4–8× faster** than GPT-6 Luna and Haiku 5.5, per Inception. (Interesting Engineering)
- **The catch:** benchmark wins are named (IFBench, GPQA Diamond, SWE-Bench Atlas) but **scores weren't published**, and **no independent tester has checked them yet**. (Startup Fortune)

Why it matters. For chat, voice and agents, **waiting is the thing users feel**. A model that is "good enough" and answers in a second can beat a smarter one that takes ten.

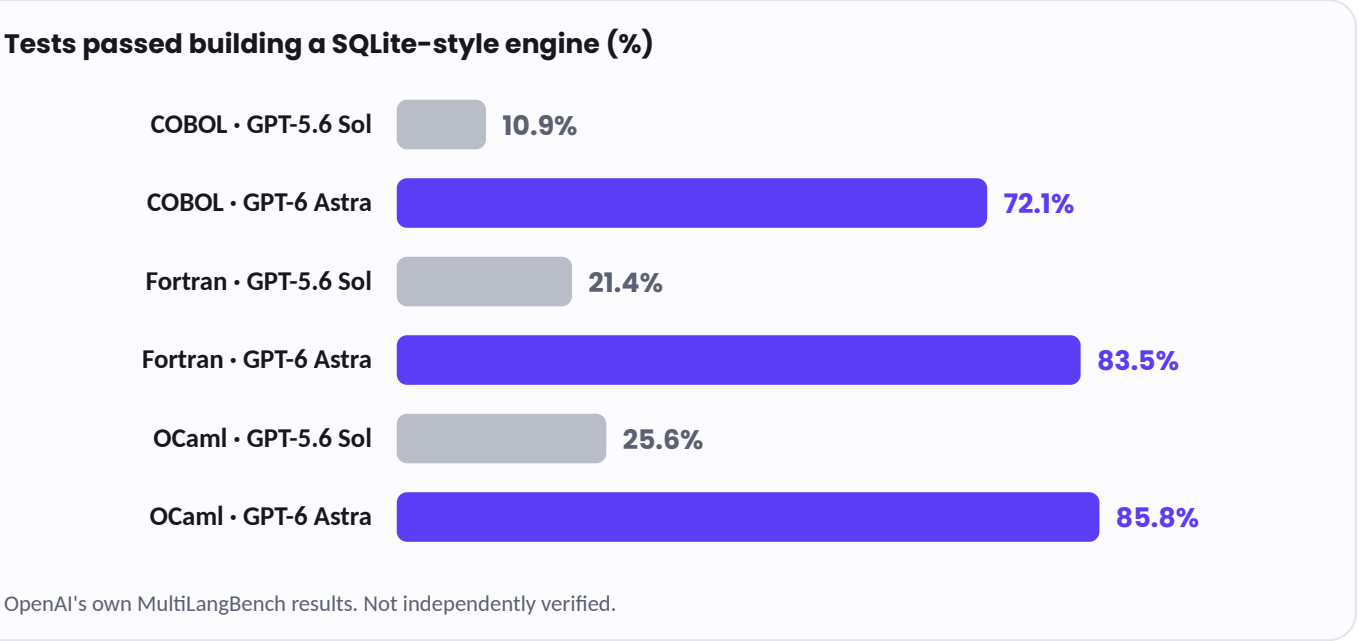
WHO SHOULD CARE 🙋 Everyone ●●○○○ 🎨 Vibe coders ●●●●○ 🛠️ Developers ●●●●○
🔥 Hype level ●●●○○

👀 **Should you care?** If your app makes **lots of small calls** (autocomplete, form-filling, agents that loop), yes: try it on OpenRouter at the launch price. **Test it on your own prompts** before trusting the benchmarks.

Top story 3 · OpenAI: GPT-6 learns COBOL

CODING LEGACY CODE

What happened. OpenAI says its newest models are much better at **"low-resource" languages**: ones with little code online to learn from, like **COBOL, Fortran, Visual Basic and OCaml**. It published a new test, **MultiLangBench**, on GitHub. (OpenAI)



OpenAI's own MultiLangBench results. Not independently verified.

The details that matter:

- **The test:** build a **SQLite-compatible database engine** in each of **20 languages**. GPT-6 Astra averages **87.4% of tests passed**, against 26.8% for GPT-5.6 Sol.
- **Migration:** moving COBOL to Java, Python, Rust and others, **GPT-6.1 Sol hit 86.6% at low effort** (vs 5.4% before) and **97.9% at higher effort**.
- **Budget option: GPT-6 Luna averages 40.3%** at about **$0.77 per attempt**, versus $3.48 for GPT-5.6 Sol.
- OpenAI itself says **real migrations still need expert review and testing**. (OpenAI)

Why it matters. Banks, insurers and governments still run on **decades-old COBOL**, and the people who know it are retiring. Mainframe migration is also **big business for Indian IT services firms**. If these numbers hold, that work gets faster, and the jobs in it change shape.

WHO SHOULD CARE 🙋 Everyone ●●○○○ 🎨 Vibe coders ●●○○○ 🛠️ Developers ●●●●○
🔥 Hype level ●●●○○

⚡ Quick hits

🇺🇸 US agencies told to rip out DeepSeek.

DHS ordered civilian agencies to remove DeepSeek products **and third-party services "developed using" them**. Experts say that phrase is **unclear for apps that route through gateways** like OpenRouter. (The Register)

🔥 PyTorch on Amazon's chips, one line.

TorchNeuron lets existing PyTorch code run on Trainium2/3 with `.to('neuron')`. Apache 2.0, public beta from 20 Oct. Deepgram reports **~60% more audio per dollar** than its GPUs. (AI Weekly)

💰 Tenstorrent raises $1.5B.

Jim Keller's chip company is now valued at **$6.2B**, more than double its December 2024 mark. Money goes into making its **Galaxy servers** at volume. (TechStartups)

🪪 A badge for your AI agent.

Sam Altman's World launched **World ID for Agents**, letting a verified human **hand proof of personhood to an agent**, as part of a bank-focused suite. Zoom's deepfake check is now generally available. (The Block)

🚫 Update: agent blocking hits a protest.

Meta says a **Claude-linked tool** browsing No Kings Coalition pages at **~4× human speed** tripped bot detection, suspending the account for **about 90 minutes**. It didn't name the tool. (AI Weekly)

☁️ Cloudflare's Clairic.

A small model for **pulling structured data out of messy text**, now on Workers AI as `@cf/cloudflare/clairic`. Cloudflare claims a pipeline using it is **3.2× faster and 49% cheaper** than LLM-only. (AI Weekly)

🏭 Inside Meta's Prometheus.

The Ohio cluster that trained **Muse Wonder** has hundreds of thousands of GPUs, aiming for **~1 million**. Meta says interruptions fell **99% from mid-2025**. (Engineering at Meta)

🗓️ Coming up.

OneDrive + Copilot event and PyTorch Conference 20 Oct. Meta earnings 21 Oct (US). MCPCon 22–23 Oct. Infosys results and Nemotron Bolt E4B 23 Oct. Mistral Large 4 weights 27 Oct.

🎯 One thing to try today: turn your best prompt into a skill

Find the prompt you've pasted most this month. Save it as `SKILL.md` with a **name** and a **"Use when…"** description, and drop it in your tool's skills folder. **Next time, just ask; the AI fetches it.** Then check your API keys aren't sitting in any shared log.

🧠 Pop quiz: 30 seconds, no Googling

1. What does an AI read first, before opening a skill?

- A The whole folder
- B Just its name and description
- C Its scripts

2. How did the researcher spot the hacker's AI agents?

- A A tip from OpenAI
- B Control panels left open online
- C A ransom note

3. How does Mercury 3 get its speed?

- A Smaller words
- B Drafting many tokens at once
- C Skipping safety checks

Answers (no peeking): 1. B: labels first (~100 tokens each); the full SKILL.md loads only when it matches · 2. B: exposed Cairn panels showed the targets, and even leaked API keys · 3. B: diffusion models refine a whole draft in parallel instead of token by token

Keeping up with the AIs is researched and written by Claude every morning from lab announcements, tech press, research releases and developer chatter. Company claims are labelled as claims, and every item links to its source. Spot an error or want a topic covered? Just reply. ✉️