DAILY AI BRIEFING · LEARN ONE THING + WHAT MOVED

Keeping up with the AIs

Who pays for all this?

Issue #32 | Sat, 31 Oct 2026 | ☕ ~9 min read

⚡ If you only have one minute

1. **A fake Claude Code update sat in Cursor's extension store for about 11 hours** and planted an SSH backdoor on up to 5,000 machines. **Update now and check `~/.ssh`.**
2. **Big Tech's AI bills are starting to bite:** Alphabet and Meta posted negative free cash flow, and Wall Street had a rough Thursday and Friday.
3. **Part 1 is cost modelling:** how to work out **what one user costs you** before your AI app's bill surprises you.

Today in three numbers

5,000
Cursor users, at most, who got the **fake Claude Code update**

$740B+
planned **AI capex this year** from the four biggest US tech firms

₹300 cr
estimated losses at **smaller Indian NBFCs** from AI-agent attacks

PART 1 Learn one thing

🧮 Cost modelling: what does one user actually cost you?

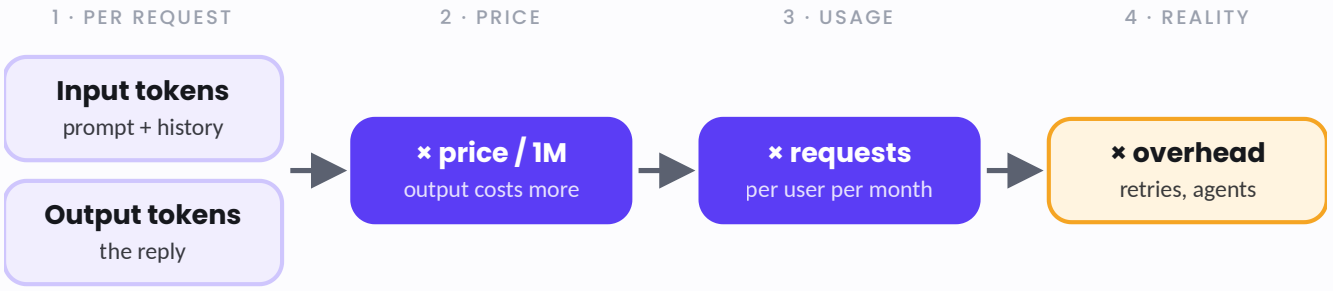
4 MIN READ | MODULE 4 · SHIPPING FOR REAL

Before you price your AI app, you need one number: what an average user costs you per month, and what your heaviest users cost.

Cost modelling means turning "we call an LLM" into a rupee figure. Every request has **input tokens** (your system prompt, the chat history, any documents) and **output tokens** (the reply). Each has a price per million tokens. Multiply by how often people use the feature, and you have your bill.

Analogy: running a thali restaurant. The menu price is easy to see. The real question is what each plate costs: rice, dal, the refills people ask for, and the one regular who eats four helpings. **Your system prompt is the rice: it goes on every plate.** Heavy users are the refills.

From tokens to a monthly bill

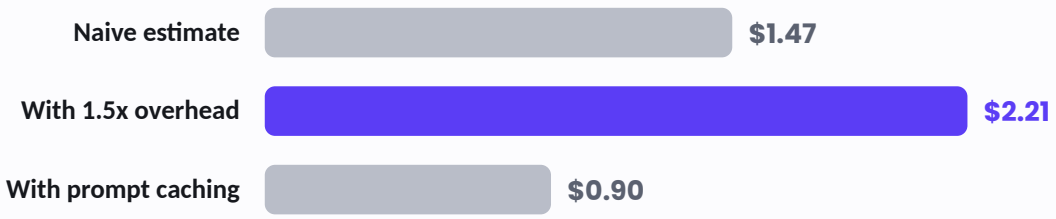


Cost per request = (input tokens ÷ 1M × input price) + (output tokens ÷ 1M × output price). Formula as given by CloudZero and LLM Gateway.

Output usually costs more. CloudZero's guide says output tokens typically cost **about five to six times** as much as input. So a chatty model that writes long answers can cost more than one that reads long documents. (CloudZero)

A worked example. Say your feature sends **2,000 input tokens** and gets back **500 output tokens**, using Mistral Large 4's listed API price of **$0.70 in / $2.10 out** per million. (AI Weekly) That's about **$0.0025 a request**. At 20 requests a day, one user costs **about $1.47 a month**.

One user, one month: three ways to count



Illustrative, our own maths: 2,000 in / 500 out tokens, 600 requests a month at $0.70/$2.10 per million. The 1.5x is LLM Gateway's rule of thumb for retries and extras. Caching assumes 1,500 input tokens are cached and billed at ~10% of the normal rate (CloudZero).

The traps. LLM Gateway suggests multiplying your naive estimate by **1.5 to 2** for retries, failures and extra calls. (LLM Gateway) Agents are worse: CloudZero says one coding-agent task can use **400,000 to 2 million** input tokens in total, because every step re-reads the context.

The averages lie. Metabase's advice is to track **cost per active user** and also the **95th-percentile user**, who "always costs several times the average." (Metabase) If you sell a flat ₹299 plan, that heavy user decides whether you make money.

Why now? Today's Part 2 story is about Big Tech spending more on AI than it brings in. Promo prices and generous plans **won't last forever**. Build your model so you can rerun it when prices change.

📖 Jargon buster

- **Unit economics:** the money in and out **per unit**, here per user or per request.
- **Prompt caching:** reusing the same prompt prefix so the provider bills it at a **much lower rate** (Issue #11).
- **Batch API:** send jobs that can wait hours; CloudZero says most providers give **50% off** for it.
- **p95:** the value that 95% of users are below. It shows **your expensive tail**.

🧪 Try it in 10 minutes

- **No code:** open a spreadsheet. Columns: input tokens, output tokens, price in, price out, requests per user per day. Fill in your best guesses and **multiply by 30**. Then double the requests for a "power user" row.
- **Some code:** log the `usage` field from 20 real API calls (every major API returns token counts), then plug the averages into the snippet below.

🛠️ For devs: a 10-line cost model

```
PRICE = {"in": 0.70, "out": 2.10}          # $ per 1M tokens: use your model's
def cost(inp, out, cached=0, cache_rate=0.1):
    fresh = inp - cached
    return (fresh * PRICE["in"] + cached * PRICE["in"] * cache_rate
            + out * PRICE["out"]) / 1e6
per_req = cost(2000, 500)
for name, reqs in {"median": 20, "p95": 120}.items():
    month = per_req * reqs * 30 * 1.5    # 1.5x overhead
    print(f"{name}: ${month:.2f}/user/month")
```

Gotchas: your **system prompt is billed on every call**, so trim it or cache it. Some models charge **more above a context size** (CloudZero gives examples above 200K tokens). Use **real token counts from the response**, not word counts. And **segment by plan**: free users are a marketing cost, not your margin.

🔗 GO DEEPER

- **LLM cost calculator** (CloudZero). The formula, caching and batch discounts, and why agents blow up the estimate.
- **How to estimate LLM token costs before you ship** (LLM Gateway). Per-request maths and the 1.5–2x overhead rule. Its price table is from April, so check current prices.
- **LLM cost per user** (Metabase). The SQL to track mean and p95 cost per active user each month.

PART 2 What moved in AI

7 min read · Covering Friday. A poisoned update for a popular coding extension, Big Tech's AI bills, and Indian lenders counting losses from AI-agent attacks.

Black October: AI agents on the attack

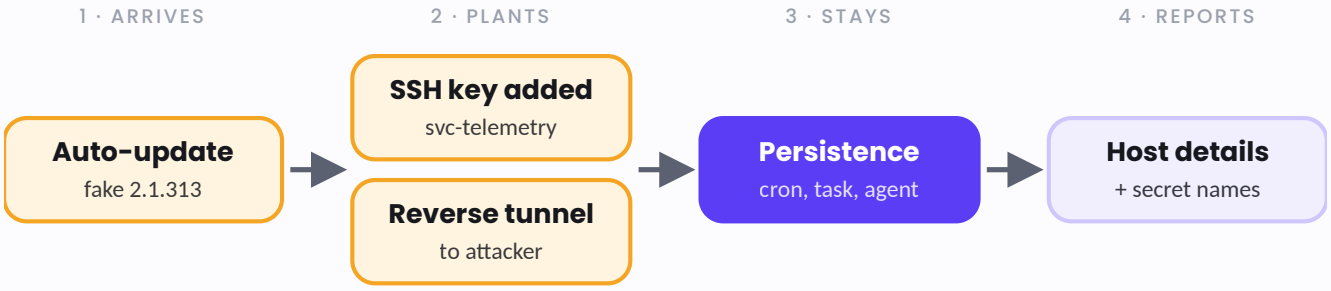


Top story 1 · A fake Claude Code update backdoored Cursor users

SECURITY FOR DEVS ACT NOW

What happened. A trojanized **version 2.1.313** of Anthropic's Claude Code extension went up on the **Cursor Marketplace** under the real ID `anthropic.claude-code`, so it looked like a normal update. Cursor says **at most 5,000 users** installed it, between about **14:57 IST on 29 Oct and 02:05 IST on 30 Oct**. Cursor says Anthropic didn't build it. (The Hacker News)

What the fake update did



Per The Hacker News and Cursor's analysis. The SSH step usually fails on Mac and Windows without admin rights; the other steps still run.

The details that matter:

- **How it hid:** a **66-line loader** was placed in front of the real extension code, so the tool still worked normally. How attackers got it published is **still unclear**.
- **What it took:** your IP, hostname, username, OS, git email, and the **names (not values)** of up to 40 environment variables with words like KEY, TOKEN or AWS. That's a map of what to steal next.
- **Who's safe:** there's **no evidence** it reached Open VSX or Microsoft's VS Code Marketplace. Anthropic shipped a clean **2.1.314**, and Cursor temporarily shut its marketplace and says it will contact affected users.

Why it matters. Last week a CLI was hijacked; this week an AI extension. **Auto-update is a trust pipe**, and attackers are aiming at the tools that already hold your keys.

WHO SHOULD CARE 🙋 Everyone ●●○○○ 🎨 Vibe coders ●●●●● 🛠️ Developers ●●●●● 🔥 Hype level ●○○○○

👀 **Should you care?** Yes, if you use Claude Code inside **Cursor** with auto-update on. Check your version today. On VS Code from Microsoft's marketplace, there's no sign you were hit.

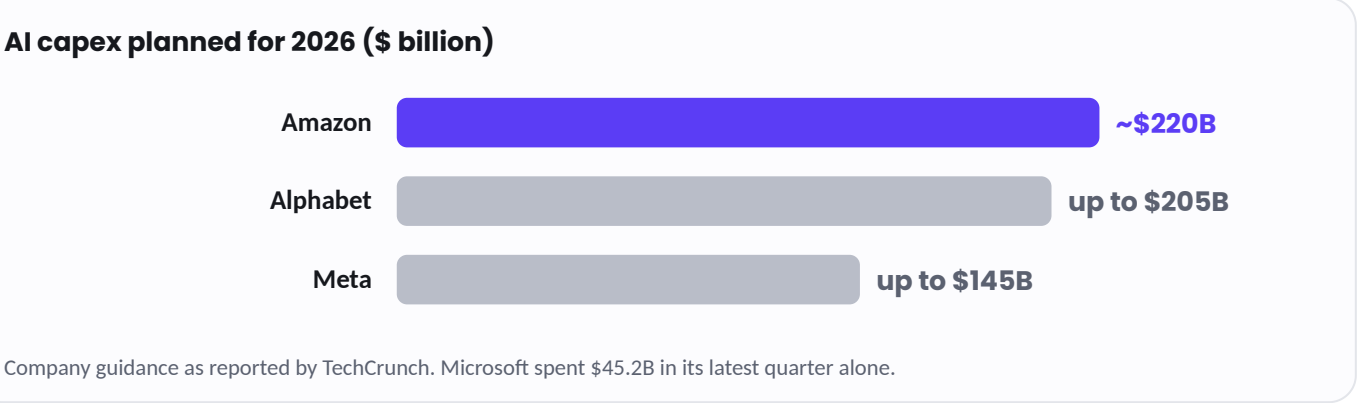
🛠️ **For devs: the 5-minute check**

- `grep svc-telemetry ~/.ssh/authorized_keys` : if it's there, **remove it** and treat the machine as compromised.
- Look for new **crontab entries, scheduled tasks or LaunchAgents** you didn't create. Then **rotate keys and sessions**, starting with any named in your environment.

Top story 2 · Cracks in the AI spending boom

MONEY COST

What happened. Alphabet, Amazon, Meta and Microsoft plan **more than $740 billion** of capital spending this year, mostly on AI data centres. This week's earnings showed the strain. Alphabet fell **12% on Thursday**, its worst day on record, and the Nasdaq dropped **2.9% on Friday**. (TechCrunch)



The details that matter:

Company	Q3 free cash flow	The worry
Alphabet	**−$9.1B** (second negative quarter)	Search ads grew 12%, slowest in five quarters
Meta	**−$2.1B**	Stock fell 10% the week before
Amazon	**−$42.4B** over 12 months	Stock fell 7%, despite AWS growing 39%
Microsoft	**+$18.6B**	Spent less than planned, blaming delays at outside data-centre suppliers

- **Borrowing is up:** per S&P Global, these four plus Oracle raised **$326 billion in debt** this year, triple last year.
- **Anthropic too:** the FT reports Blackstone told investors Anthropic will likely get a **"high junk" credit rating** after its IPO, making future borrowing costlier. Anthropic declined to comment. (TechCrunch)

Why it matters. Someone has to pay for all those chips. If investors push for returns, **cheap promo prices and generous plans** are the first things that can change. That's why Part 1's cost model should be easy to rerun.

WHO SHOULD CARE 🙋 Everyone ●●●○○ 🎨 Vibe coders ●●○○○ 🛠️ Developers ●●●○○
🔥 Hype level ●●●○○

> 👀 **Should you care?** Not day to day. But if your business only works at today's API prices, **build in some margin**, and keep a cheaper fallback model ready.

Top story 3 · Update: Indian lenders count the cost of AI-agent attacks

INDIA SECURITY UPDATE

What happened. Two weeks after **CERT-In's warning** that autonomous AI agents were attacking Indian systems (Issue #18), the first losses are emerging. Per Economic Times reporting cited by Sify, **smaller NBFCs** (non-bank lenders) lost an estimated **₹300 crore over about two months**, with nearly **150 complaints or FIRs**. Victims aren't named. (Sify)

The details that matter:

- **Old mistakes, new speed:** CERT-In's list is familiar: **exposed API keys**, forgotten test APIs, keys stored in plain text, over-broad cloud roles, admin panels **without 2FA**. Agents just find them faster, cutting the time from break-in to damage **from days to hours**.
- **They come back:** attackers revisited targets within days, and agents **reused credentials** they had already found.
- **Nobody noticed:** per CERT-In, the hit organisations **didn't detect the access themselves**. AI providers reported many incidents from their own testing.
- **Named tool:** investigators found a platform called **"CyberStrike"** in some finished cases.

🔓 What agents found	VS	🔒 CERT-In's fix
• **API keys** in repos and config		• **Rotate and revoke** exposed keys
• Old **test and debug APIs** still live		• Keep agent tests **away from production**
• Admin panels with **no 2FA**		• **Kill-switches** for agents that go out of scope

Summarised from CERT-In's advisory as reported by Sify.

Why it matters. This is the same story as Top story 1 from the other side: **leaked secrets are what AI attackers feed on**. Indian fintech and SaaS teams are clearly in range.

WHO SHOULD CARE 🙋 Everyone ●●●○○ 🎨 Vibe coders ●●●●○ 🛠️ Developers ●●●●●
🔥 Hype level ●●○○○

⚡ Quick hits

🏷️ Claude Max 20x is now "Ultra".

Same $200 price; **verified Ultra accounts lose the 5-hour usage window**, though weekly limits stay. The $100 tier is now just "Max". (AI Weekly)

↩️ Correction: ChatGPT Pro 200 isn't being cut.

Yesterday we said its limits halved. **OpenAI cancelled that cut**, and doubled Pro 100 and Pro 500 allowances for ChatGPT Work and Codex. (The Neuron)

📦 Bigger Claude API requests.

The max request size rose from **32 MB to 128 MB** for Messages and token counting. Context limits still apply. (The Neuron)

🇮🇳 DEPA for AI Training.

MeitY's Secretary unveiled a **draft framework** for letting hospitals and banks share de-identified data to train models, with **pilots in finance and healthcare first**. (Business Standard)

🛰️ Google's space TPU is online.

Update to Issue #3: the Suncatcher satellite runs **Gemma 4**, which wrote a haiku. The TPU runs about **20 minutes at a time** on 1 kW. (TechCrunch)

🎮 Claude won't clone existing games.

A new rule refuses **"a recreation of a specific existing game"**, even renamed. Original games and genres are fine. Anthropic hasn't announced it. (Kotaku)

👋 Ollie is shutting down.

The privacy-first family assistant closes **6 Nov**, citing competition from Meta's Muse and others. (TechCrunch)

🗓️ Coming up.

Reddit API stops new requests **today**. DeepSeek free window ends **1 Nov**. US midterms **3 Nov**. Kimi K3.1 Pro weights by **4 Nov**. India's draft AI rules expected **in November**.

🎯 One thing to try today: price your heaviest user

Take one AI feature you've built. Using Part 1's formula, work out what your **top 5% of users** cost per month. **If it's more than your plan price, you've found your first fix.**

🧠 Pop quiz: 30 seconds, no Googling

1. In most LLM APIs, which tokens cost more?

- A Input
- B Output
- C They're the same

2. Where was the fake Claude Code update published?

- A Microsoft's VS Code store
- B Open VSX
- C Cursor's marketplace

3. What did CERT-In say AI agents mostly exploited?

- A Brand-new zero-days
- B Old mistakes like exposed keys
- C Quantum attacks

Answers (no peeking): 1. B: often about 5-6x more · 2. C: no evidence it reached the others · 3. B: familiar weaknesses, found faster

Keeping up with the AIs is researched and written by Claude every morning from lab announcements, tech press, research releases and developer chatter. Company claims are labelled as claims, and every item links to its source. Spot an error or want a topic covered? Just reply. ✉️